

LevelDB Study

Week 5

2022. 07. 05

Presented by Min-guk Choi

koreachoi96@gmail.com

1. Experiment Report

2. How to analyze LevelDB

- Experiment Report
 - ✓ leveldb-wiki (kor)
 - ✓ Notice

leveldb-wiki (kor)

☰ README.md

leveldb-wiki (kor)

Background

1. What is a key-value store?

2. Why open source?

3. Job market

Analysis

0. Overall

1. Key-Value Interface

2. WAL

3. Memtable

4. Compaction

5. SSTable

6. Bloom Filter

7. Cache

8. Manifest

Benchmarks Experiments

• WAL

• Memtable


• Compaction

• SSTable







• Bloom Filter


• Cache

🔗 main ▾ leveldb-wiki / benchmarks /

korea-choi readme update

..

 bloomfilter.md	readme update
 cache.md	readme update
 compaction.md	readme update
 memtable.md	readme update
 sstable.md	readme update
 wal.md	readme update



<https://github.com/DKU-StarLab/leveldb-wiki>

☰ 11 lines (6 sloc) | 96 Bytes

WAL

Hypothesis

Design

Experiment Enviornment

Result

Discussion

Notice

- Write your experiment report into markdown documents.
 - Don't forget to cite the source of the referenced figures.
 - Don't forget to add a watermark on your figures and graphs.
- Upload your documents through pull-request.
 - You can also change front read page if you need.
 - Pull request until **Tuesday, 8/2 11AM.**
 - <https://github.com/DKU-StarLab/leveldb-wiki>
- It will be finally uploaded at <https://sslabor.dankook.ac.kr/leveldb-wiki>

1. Experiment Report

2. How to analyze LevelDB




- How to analyze LevelDB
 - ✓ What/When to use
 - ✓ Where to start?
 - ✓ Source Code
 - ✓ Code Flow
 - ✓ Watermark
 - ✓ PPT

What/When to use?

	1	2
Static Analysis Tool - Remarks & Code	VS code - Go to Definition - Go to References - Search - File Explorer	Understand - Class Diagram - Object References - Butter Fly
Dynamic Analysis Tool - Code Flow/Tracing	Uftrace - Record - Replay - Tracing	GDB - Break point - Step into / Next - Print / Display - Line by Line - Argument, Variable

Where to start?

- Start with basic operations
 - Put(Write) operations
 - `db_bench --benchmarks="fillseq, fillrandom"`
 - `db_bench.cc:DoWrite`
 - Get(Read) operations
 - `db_bench --benchmarks="readrandom"`
 - `db_bench.cc:ReadRandom`
 - Seek(Scan) operations
 - `db_bench --benchmarks="readseq"`
 - `db_bench.cc:ReadSequential`



Source Code

Topics

Topic	User API	Internal Operation	Contents	Source Code
WAL/Manifest	Put	Open Flush Compaction	Version Control Log Format Manifest Format	log_.h version_.h db_impl.h repair.cc
Memtable	Put Get Iterator	Flush	Skiplist Arena Batch Write	wirte_batch_internal.h skiplist.h memtable.h db_impl.h arena.h
Compaction	Put	Compaction	Compaction Policy Merge Iterator	db_impl.h merger.h version_set.h
SSTable	Get	Flush	SST Format	table/

<https://github.com/DKU-StarLab/leveldb-study/blob/main/analysis/README.md>

Function

WAL/Manifest	Memtable
<code>leveldb::log::Writer::AddRecord</code> <code>leveldb::_GLOBAL__N_1::PosixWritableFile::Append</code> <code>leveldb::log::Writer::EmitPhysicalRecord</code> <code>leveldb::VersionSet::PickCompaction</code> <code>leveldb::VersionSet::LogAndApply</code> <code>leveldb::VersionSet::AppendVersion</code> <code>leveldb::VersionSet::NeedsCompaction</code> <code>leveldb::VersionEdit::VersionEdit</code> <code>leveldb::Version::OverlapInLevel</code> <code>leveldb::DBImpl::MaybeScheduleCompaction</code> <code>leveldb::DBImpl::InstallCompactionResults</code> <code>leveldb::_GLOBAL__N_1::PosixWritableFile::SyncFd</code> <code>leveldb::_GLOBAL__N_1::PosixEnv::RemoveFile</code>	<code>leveldb::Version::PickLevelForMemTableOutput</code> <code>leveldb::DBImpl::CompactMemTable</code> <code>leveldb::DBImpl::WriteLevel0Table</code> <code>leveldb::SkipList::RandomHeight</code> <code>leveldb::SkipList::NewNode</code> <code>leveldb::SkipList::Equal</code> <code>leveldb::SkipList::Node::SetNext</code> <code>leveldb::SkipList::Iterator</code> <code>leveldb::Arena::MemoryUsage</code> <code>leveldb::Arena::AllocateNewBlock</code> <code>leveldb::Arena::~Arena</code>

Function

Compaction	SSTable
<ul style="list-style-type: none">leveldb::DBImpl::BGWorkleveldb::DBImpl::DoCompactionWorkleveldb::DBImpl::FinishCompactionOutputFileleveldb::DBImpl::InstallCompactionResultsleveldb::DBImpl::CleanupCompactionleveldb::NewTwoLevelIteratorleveldb::NewMergingIteratorleveldb::_GLOBAL__N_1::MergingIterator::SeekToFirstleveldb::_GLOBAL__N_1::MergingIterator::key	<ul style="list-style-type: none">leveldb::ReadBlockleveldb::TableBuilder::Addleveldb::TableBuilder::Flushleveldb::TableBuilder::WriteBlockleveldb::port::Snappy_Compressleveldb::_GLOBAL__N_1::HandleTable::Lookupleveldb::_GLOBAL__N_1::PosixMmapReadableFile::Readleveldb::_GLOBAL__N_1::TwoLevelIterator::InitDataBlockleveldb::_GLOBAL__N_1::PosixWritableFile::Appendleveldb::_GLOBAL__N_1::PosixWritableFile::Sync

Function

Bloom Filter	Cache
<ul style="list-style-type: none">leveldb::Table::ReadFilterleveldb::FilterBlockBuilder::Finishleveldb::FilterBlockBuilder::StartBlockleveldb::FilterBlockBuilder::GenerateFilterleveldb::FilterBlockBuilder::FilterBlockBuilderleveldb::FilterBlockBuilder::AddKeyleveldb::FilterBlockReader::KeyMayMatchleveldb::_GLOBAL__N_1::BloomFilterPolicy::KeyMayMatchleveldb::FilterBlockReader::FilterBlockReaderleveldb::_GLOBAL__N_1::BloomHashleveldb::_GLOBAL__N_1::BloomFilterPolicy::CreateFilterleveldb::_GLOBAL__N_1::BloomFilterPolicy::Name	<ul style="list-style-type: none">leveldb::TableCache::FindTableleveldb::DeleteCachedBlockleveldb::_GLOBAL__N_1::LRUCache::Releaseleveldb::_GLOBAL__N_1::LRUCache::LRU_Removeleveldb::_GLOBAL__N_1::LRUCache::LRU_Appendleveldb::_GLOBAL__N_1::LRUCache::FinishEraseleveldb::_GLOBAL__N_1::LRUCache::Unrefleveldb::_GLOBAL__N_1::ShardedLRUCache::Shardleveldb::_GLOBAL__N_1::ShardedLRUCache::Lookupleveldb::_GLOBAL__N_1::ShardedLRUCache::Insertleveldb::_GLOBAL__N_1::PosixMmapReadableFile::Read

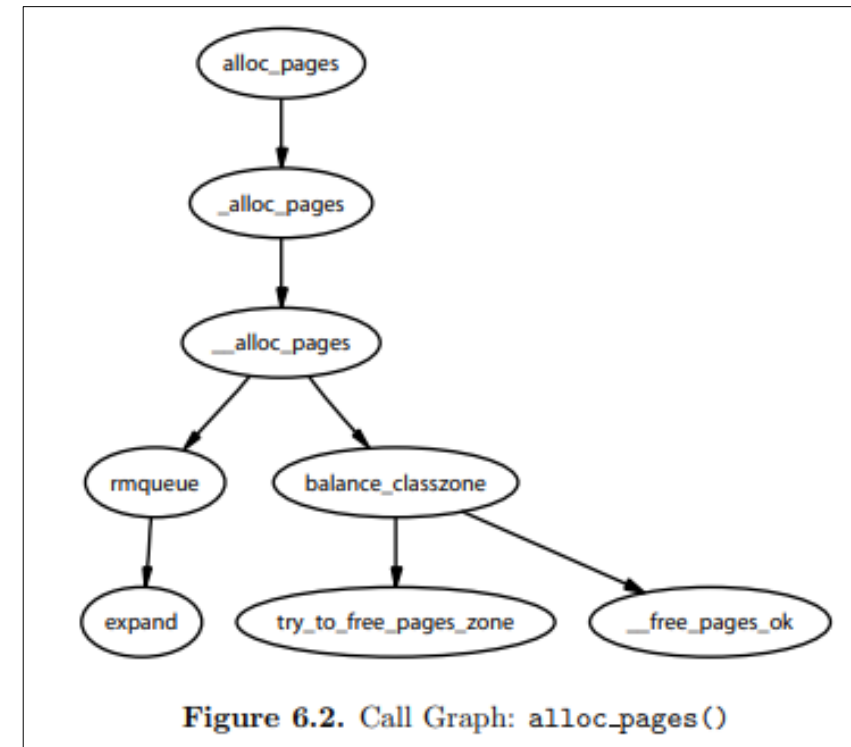
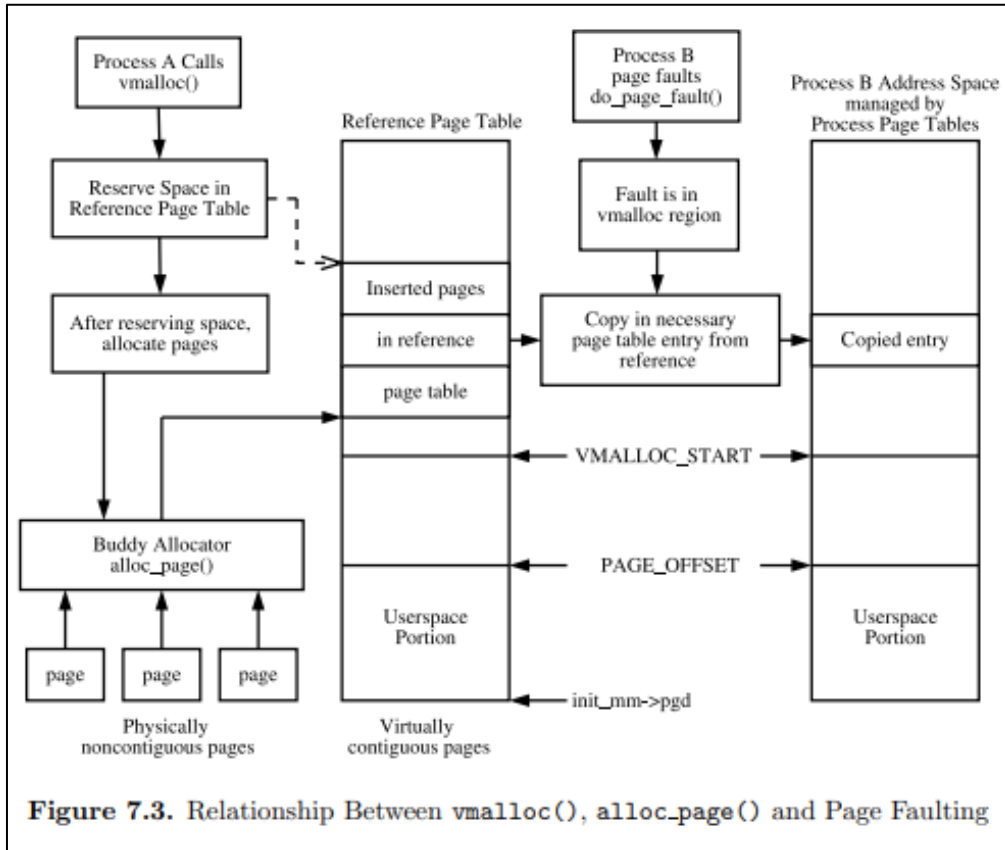
Code Flow

```
===== Back-trace =====
5.188 s : (100000) leveldb::SkipList::Insert
5.188 s : (100000) leveldb::MemTable::Add
5.188 s : (100000) leveldb::_GLOBAL__N_1::MemTableInserter::Put
5.188 s : (100000) leveldb::WriteBatch::Iterate
5.188 s : (100000) leveldb::WriteBatchInternal::InsertInto
5.188 s : (100000) leveldb::DBImpl::Write
5.188 s : (100000) leveldb::Benchmark::DoWrite
5.188 s : (100000) leveldb::Benchmark::WriteRandom
5.188 s : (100000) leveldb::Benchmark::ThreadBody
:
===== Call Graph =====
5.188 s : (100000) leveldb::SkipList::Insert
4.824 s : |-(100000) leveldb::SkipList::FindGreaterOrEqual
13.450 ms : |-(100000) leveldb::SkipList::GetMaxHeight
:
473.144 ms : |-(2495692) leveldb::SkipList::Node::Next
:
4.121 s : |-(2495692) leveldb::SkipList::KeyIsAfterNode
3.903 s : |-(2404364) leveldb::MemTable::KeyComparator::operator()
993.576 ms : |-(4808728) leveldb::GetLengthPrefixedSlice
:
2.544 s : |-(2404364) leveldb::InternalKeyComparator::Compare
492.141 ms : |-(2404364) leveldb::_GLOBAL__N_1::BytewiseComparatorImpl::Compare
:
25.035 ms : |-(100000) leveldb::SkipList::RandomHeight
16.364 ms : |-(133759) leveldb::Random::OneIn
4.255 ms : |-(133759) leveldb::Random::Next
:
11.891 ms : |-(100026) leveldb::SkipList::GetMaxHeight
:
30.790 ms : |-(100000) leveldb::SkipList::NewNode
5.508 ms : |-(100000) leveldb::Arena::AllocateAligned
1.455 ms : |-(890) leveldb::Arena::AllocateFallback
1.384 ms : |-(890) leveldb::Arena::AllocateNewBlock
```

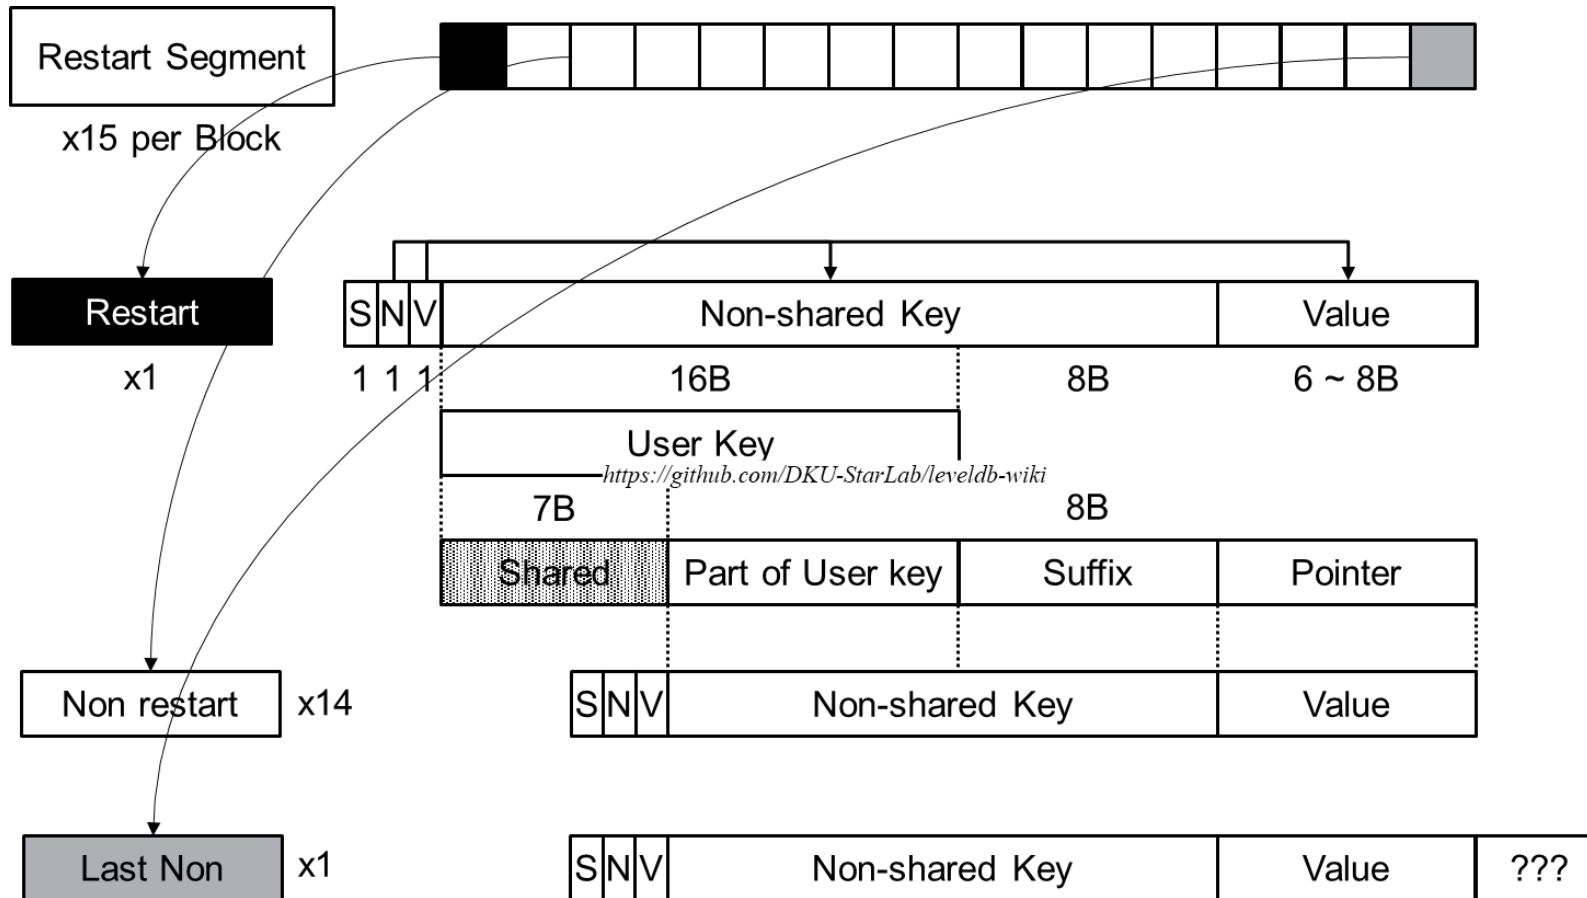
Help: (press any key to exit)

ARROW	Navigation
PgUp/Dn	
Home/End	
Enter	Fold/unfold graph or Select session
G	Show (full) call graph
g	Show call graph for this function
R	Show uftrace report
r	Show uftrace report for this function
s	Sort by the next column in report
I	Show uftrace info
S	Change session
O	Open editor
c/e	Collapse/Expand direct children graph
C/E	Collapse/Expand all descendant graph
n/p	Next/Prev sibling
u	Move up to parent
l	Move to the longest executed child
j/k	Move down/up
z	Set current line to the center of screen
/	Search
</>/N/P	Search next/prev
v	Show debug message
f	Customize fields in graph or report mode
h/?	Show this help
q	Quit

Example Figure

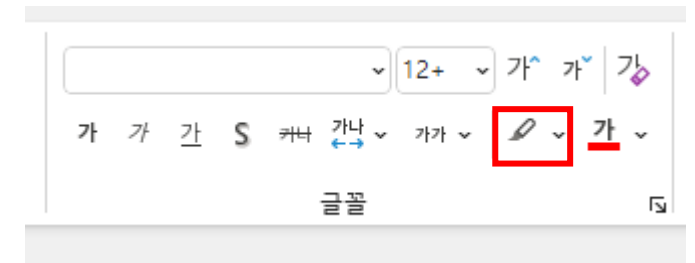


Watermark




Watermark

- <https://github.com/DKU-StarLab/leveldb-wiki>
- Use powerpoint or other programs/websites.
- Add a background color on the watermark.





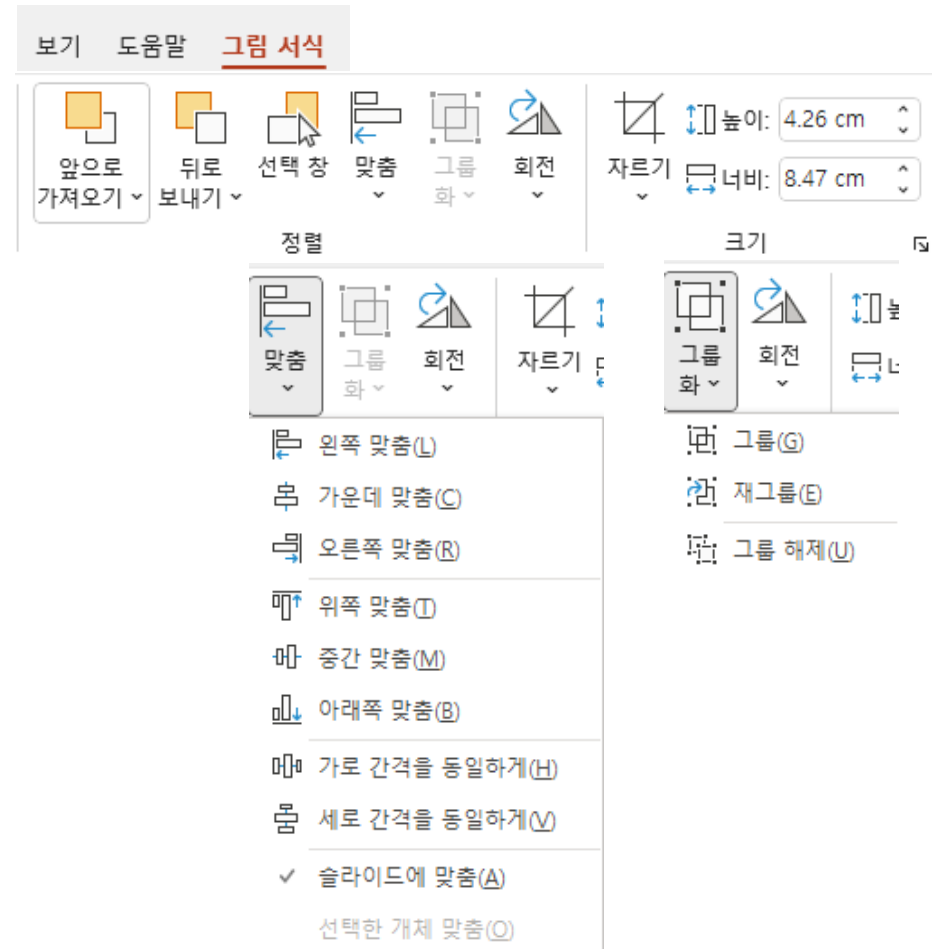
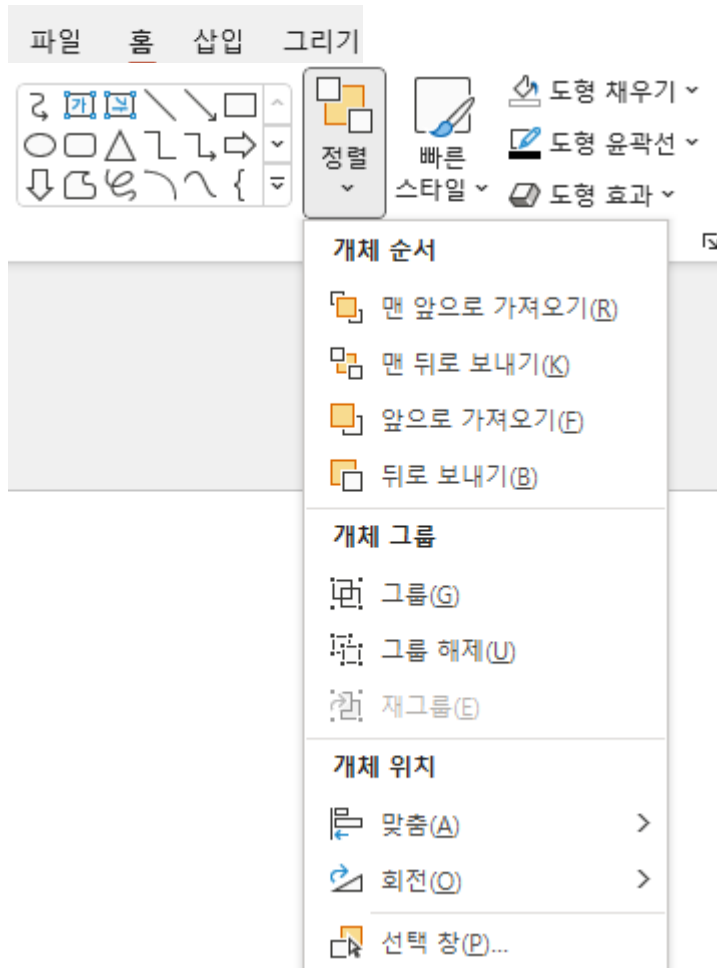
그것은 무한한 시간과 노가다의 집합체였다



대학원 탈출일지 12화 - 피피티

<https://comic.naver.com/webtoon/detail?titleId=790713&no=13&weekday=tue>

PPT



Thank you