

# LevelDB Study

## Bloom Filter Analysis

Made by Kim Han Su

E-Mail: [khs20010327@naver.com](mailto:khs20010327@naver.com)

# UFTRACE

---

```
[ 1418] | leveldb::Benchmark::Benchmark() {  
[ 1418] | leveldb::NewBloomFilterPolicy() {  
0.412 us [ 1418] | operator new();  
[ 1418] | leveldb::_GLOBAL__N_1::BloomFilterPolicy::BloomFilterPolicy() {  
0.052 us [ 1418] | leveldb::FilterPolicy::FilterPolicy();  
0.288 us [ 1418] | } /* leveldb::_GLOBAL__N_1::BloomFilterPolicy::BloomFilterPolicy */  
1.455 us [ 1418] | } /* leveldb::NewBloomFilterPolicy */  
0.043 us [ 1418] | leveldb::WriteOptions::WriteOptions();  
[ 1418] | leveldb::BytewiseComparator() {  
0.043 us [ 1418] | leveldb::NoDestructor::get();
```

# NewBloomFilterPolicy

---

```
const FilterPolicy* NewBloomFilterPolicy(int bits_per_key) {  
    return new BloomFilterPolicy(bits_per_key);  
}
```

```
[ 1418] | leveldb::Benchmark::Benchmark() {  
[ 1418] | leveldb::NewBloomFilterPolicy() {  
0.412 us [ 1418] | operator new();  
[ 1418] | leveldb::_GLOBAL__N_1::BloomFilterPolicy::BloomFilterPolicy() {
```

# FilterPolicy

---

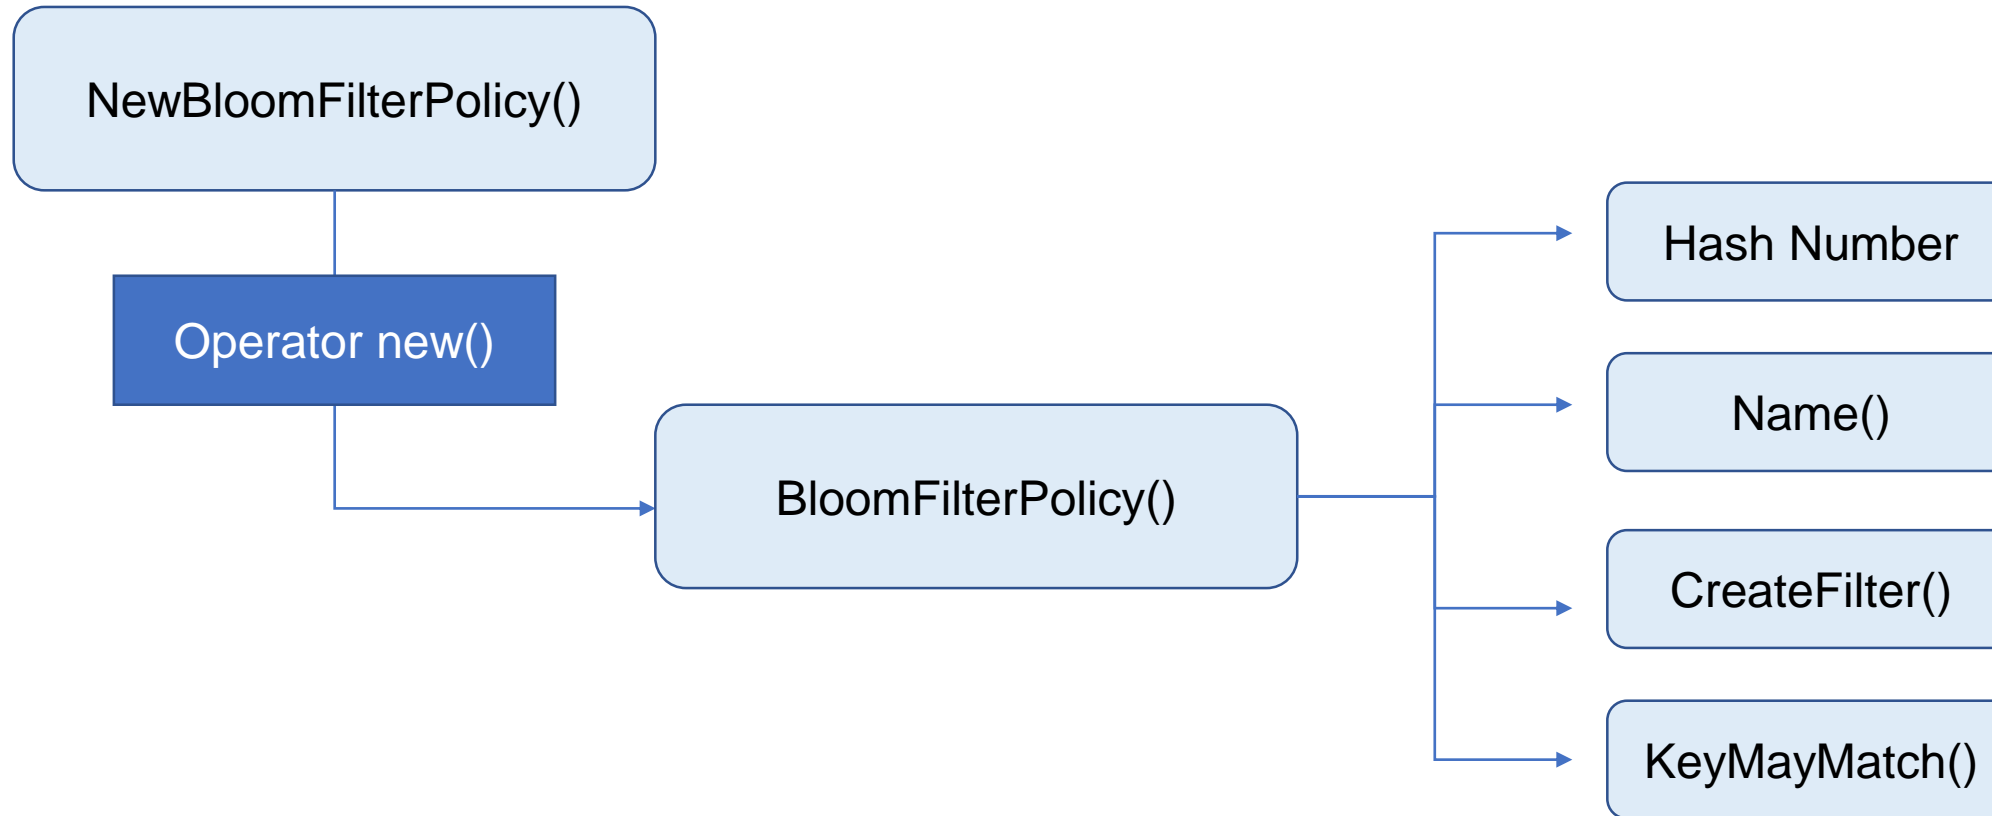
```
class LEVELDB_EXPORT FilterPolicy {
public:
    virtual ~FilterPolicy();

    // Return the name of this policy. Note that if the filter encoding
    // changes in an incompatible way, the name returned by this method
    // must be changed. Otherwise, old incompatible filters may be
    // passed to methods of this type.
    virtual const char* Name() const = 0;

    virtual void CreateFilter(const Slice* keys, int n,
                             std::string* dst) const = 0;

    virtual bool KeyMayMatch(const Slice& key, const Slice& filter) const = 0;
};
```

# CodeFlow



# Hash Number & Name()

---

```
class BloomFilterPolicy : public FilterPolicy {
public:
    explicit BloomFilterPolicy(int bits_per_key) : bits_per_key_(bits_per_key) {
        // We intentionally round down to reduce probing cost a little bit
        k_ = static_cast<size_t>(bits_per_key * 0.69); // 0.69 ≈ ln(2)
        if (k_ < 1) k_ = 1;
        if (k_ > 30) k_ = 30;
    }

    const char* Name() const override { return "leveldb.BuiltinBloomFilter2"; }
};
```

# CreateFilter

---

```
void CreateFilter(const Slice* keys, int n, std::string* dst) const override {  
    // Compute bloom filter size (in both bits and bytes)  
    size_t bits = n * bits_per_key_;  
  
    // For small n, we can see a very high false positive rate.  Fix it  
    // by enforcing a minimum bloom filter length.  
    if (bits < 64) bits = 64;  
  
    size_t bytes = (bits + 7) / 8;  
    bits = bytes * 8;
```

# CreateFilter

---

```
const size_t init_size = dst->size();
dst->resize(init_size + bytes, 0);
dst->push_back(static_cast<char>(k_)); // Remember # of probes in filter
char* array = &(*dst)[init_size];
for (int i = 0; i < n; i++) {
    // Use double-hashing to generate a sequence of hash values.
    // See analysis in [Kirsch,Mitzenmacher 2006].
    uint32_t h = BloomHash(keys[i]);
    const uint32_t delta = (h >> 17) | (h << 15); // Rotate right 17 bits
    for (size_t j = 0; j < k_; j++) {
        const uint32_t bitpos = h % bits;
        array[bitpos / 8] |= (1 << (bitpos % 8));
        h += delta;
    }
}
```



# CreateFilter

---

```
const size_t init_size = dst->size();
dst->resize(init_size + bytes, 0);
dst->push_back(static_cast<char>(k_)); // Remember # of probes in filter
char* array = &(*dst)[init_size];
for (int i = 0; i < n; i++) {
    // Use double-hashing to generate a sequence of hash values.
    // See analysis in [Kirsch,Mitzenmacher 2006].
    uint32_t h = BloomHash(keys[i]);
    const uint32_t delta = (h >> 17) | (h << 15); // Rotate right 17 bits
    for (size_t j = 0; j < k_; j++) {
        const uint32_t bitpos = h % bits;
        array[bitpos / 8] |= (1 << (bitpos % 8));
        h += delta;
    }
}
}
```

# BloomHash & Hash

---

```
namespace {  
static uint32_t BloomHash(const Slice& key) {  
    return Hash(key.data(), key.size(), 0xbc9f1d34);  
}
```

```
uint32_t Hash(const char* data, size_t n, uint32_t seed) {  
    // Similar to murmur hash  
    const uint32_t m = 0xc6a4a793;  
    const uint32_t r = 24;  
    const char* limit = data + n;  
    uint32_t h = seed ^ (n * m);
```

# CreateFilter

---

```
const size_t init_size = dst->size();
dst->resize(init_size + bytes, 0);
dst->push_back(static_cast<char>(k_)); // Remember # of probes in filter
char* array = &(*dst)[init_size];
for (int i = 0; i < n; i++) {
    // Use double-hashing to generate a sequence of hash values.
    // See analysis in [Kirsch,Mitzenmacher 2006].
    uint32_t h = BloomHash(keys[i]);
    const uint32_t delta = (h >> 17) | (h << 15); // Rotate right 17 bits
    for (size_t j = 0; j < k_; j++) {
        const uint32_t bitpos = h % bits;
        array[bitpos / 8] |= (1 << (bitpos % 8));
        h += delta;
    }
}
```

# KeyMayMatch

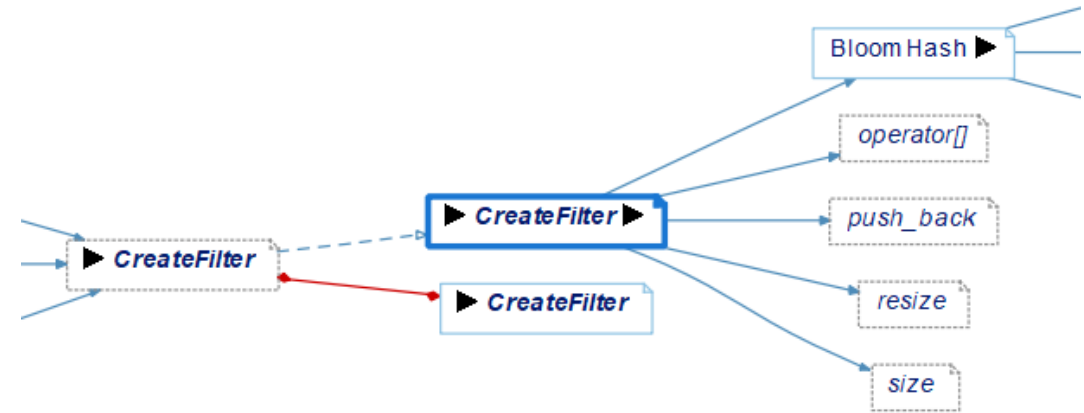
---

```
uint32_t h = BloomHash(key);  
const uint32_t delta = (h >> 17) | (h << 15); // Rotate right 17 bits  
for (size_t j = 0; j < k; j++) {  
    const uint32_t bitpos = h % bits;  
    if ((array[bitpos / 8] & (1 << (bitpos % 8))) == 0) return false;  
    h += delta;  
}  
return true;  
}
```

# InternalFilter?

[ 1471] |  
0.039 us [ 1471] |  
[ 1471] |  
0.041 us [ 1471] |  
0.152 us [ 1471] |  
[ 1471] |  
0.034 us [ 1471] |  
0.143 us [ 1471] |  
[ 1471] |  
0.052 us [ 1471] |

```
leveldb::DBImpl::DBImpl() {  
    leveldb::DB::DB();  
    leveldb::InternalKeyComparator::InternalKeyComparator() {  
        leveldb::Comparator::Comparator();  
    } /* leveldb::InternalKeyComparator::InternalKeyComparator */  
    leveldb::InternalFilterPolicy::InternalFilterPolicy() {  
        leveldb::FilterPolicy::FilterPolicy();  
    } /* leveldb::InternalFilterPolicy::InternalFilterPolicy */  
    leveldb::SanitizeOptions() {  
        leveldb::ClipToRange();  
    }  
}
```



# InternalFilterPolicy

---

```
// Filter policy wrapper that converts from internal keys to user keys
class InternalFilterPolicy : public FilterPolicy {
private:
    const FilterPolicy* const user_policy_;

public:
    explicit InternalFilterPolicy(const FilterPolicy* p) : user_policy_(p) {}
    const char* Name() const override;
    void CreateFilter(const Slice* keys, int n, std::string* dst) const override;
    bool KeyMayMatch(const Slice& key, const Slice& filter) const override;
};
```

# InternalFilterPolicy

---

```
void InternalFilterPolicy::CreateFilter(const Slice* keys, int n,
                                       std::string* dst) const {
    // We rely on the fact that the code in table.cc does not mind us
    // adjusting keys[].
    Slice* mkey = const_cast<Slice*>(keys);
    for (int i = 0; i < n; i++) {
        mkey[i] = ExtractUserKey(keys[i]);
        // TODO(sanjay): Suppress dups?
    }
    user_policy_ -> CreateFilter(keys, n, dst);
}
```

# ExtractUserKey

---

```
// Returns the user key portion of an internal key.  
inline Slice ExtractUserKey(const Slice& internal_key) {  
    assert(internal_key.size() >= 8);  
    return Slice(internal_key.data(), internal_key.size() - 8);  
}
```



# ExtractUserKey

---

```
// Returns the user key portion of an internal key.  
inline Slice ExtractUserKey(const Slice& internal_key) {  
    assert(internal_key.size() >= 8);  
    return Slice(internal_key.data(), internal_key.size() - 8);  
}
```

# ExtractUserKey

---

```
// Returns the user key portion of an internal key.
inline Slice ExtractUserKey(const Slice& internal_key) {
    assert(internal_key.size() >= 8);
    return Slice(internal_key.data(), internal_key.size() - 8);
}
```

## Google's original internal key

Google's original leveldb internal keys comprise of four components:

- total internal key size
- user's binary key
- 7 byte sequence number
- 1 byte type code

The internal key size includes the user's binary key, sequence number, and type code.

- Key
  - userkey: passed from user, in Slice format
  - InternalParsedKey: userkey + seqNum + valuetype

# InternalFilterPolicy

---

```
const char* InternalFilterPolicy::Name() const { return user_policy_->Name(); }

void InternalFilterPolicy::CreateFilter(const Slice* keys, int n,
                                       std::string* dst) const {
    // We rely on the fact that the code in table.cc does not mind us
    // adjusting keys[].
    Slice* mkey = const_cast<Slice*>(keys);
    for (int i = 0; i < n; i++) {
        mkey[i] = ExtractUserKey(keys[i]);
        // TODO(sanjay): Suppress dups?
    }
    user_policy_->CreateFilter(keys, n, dst);
}

bool InternalFilterPolicy::KeyMayMatch(const Slice& key, const Slice& f) const {
    return user_policy_->KeyMayMatch(ExtractUserKey(key), f);
}
```

# QnA

---

