

# LevelDB Study

## Bloom Filter Analysis

Made by Kim Han Su

E-Mail: [khs20010327@naver.com](mailto:khs20010327@naver.com)

# Contents

---

- More details about last presentation
- About “ReadMissing”


# Hypothesis

---

- LevelDB는  $\ln 2 * \text{bits\_per\_key}$ 개의 해시 함수를 사용
- -> Bits per key 값이 10이라면 이론상 6.9개의 해시 함수를 사용해야함
- -> 실제론 6개를 사용
- $9 \rightarrow 9 \times 0.69 = 6.21 \rightarrow 6$
- $12 \rightarrow 12 \times 0.69 = 8.28 \rightarrow 8$

# Result

latency




# Code edit

---

```
class BloomFilterPolicy : public FilterPolicy {
public:
    explicit BloomFilterPolicy(int bits_per_key) : bits_per_key_(bits_per_key) {
        // We intentionally round down to reduce probing cost a little bit
        k_ = static_cast<size_t>(bits_per_key * 0.71); // 0.69 ≈ ln(2)
        if (k_ < 1) k_ = 1;
        if (k_ > 30) k_ = 30;
    }
};
```

# Result



# Mathematical calculation

---

$$(1 - e^{-\ln 2})^{\ln 2 \cdot 10}$$

소수 형태:

0.00819254...

$$\left(1 - e^{\frac{-7}{10}}\right)^7$$

소수 형태:

0.00819372...

$$\left(1 - e^{\frac{-6}{10}}\right)^6$$

소수 형태:

0.00843620...

# Mathematical calculation

---

$$\left(1 - e^{\frac{-7}{10}}\right)^7$$

소수 형태:

0.00819372...

$$\left(1 - e^{\frac{-6}{10}}\right)^6$$


소수 형태:

0.00843620...

- 0.00024 => 240 per 1 Million



# Result




# Mathematical calculation

$$\left(1 - e^{\frac{-5}{10}}\right)^5$$

소수 형태:  
0.00943092 ...


$$\left(1 - e^{\frac{-8}{10}}\right)^8$$

소수 형태:  
0.00845547 ...



# Result

latency



# Mathematical calculation

---

$$\left(1 - e^{\frac{-8}{12}}\right)^8$$

소수 형태:  
0.00314235 ...

$$\left(1 - e^{\frac{-6}{10}}\right)^6$$

소수 형태:  
0.00843620 ...

$$\left(1 - e^{\frac{-6}{9}}\right)^6$$

소수 형태:  
0.01327213 ...


# Result

$$\left(1 - e^{\frac{-8}{12}}\right)^8$$

소수 형태:  
0.00314235...

$$\left(1 - e^{\frac{-6}{10}}\right)^6$$

소수 형태:  
0.00843620...



# Mathematical calculation

$$\left(1 - e^{\frac{-6}{10}}\right)^6$$

소수 형태:

0.00843620...


$$\left(1 - e^{\frac{-6}{9}}\right)^6$$

소수 형태:

0.01327213...

$$y = \left(1 - e^{\left(-\frac{6}{x}\right)}\right)^6$$

latency



# Conclusion

---

- 1. 같은 hash면 bloom\_bits값이 큰 쪽이 읽기 성능이 좋다
- 2. hash 개수에 비해 bloom\_bits가 false positive에 영향이 크다

$$\left(1 - e^{-\frac{6}{10}}\right)^6$$

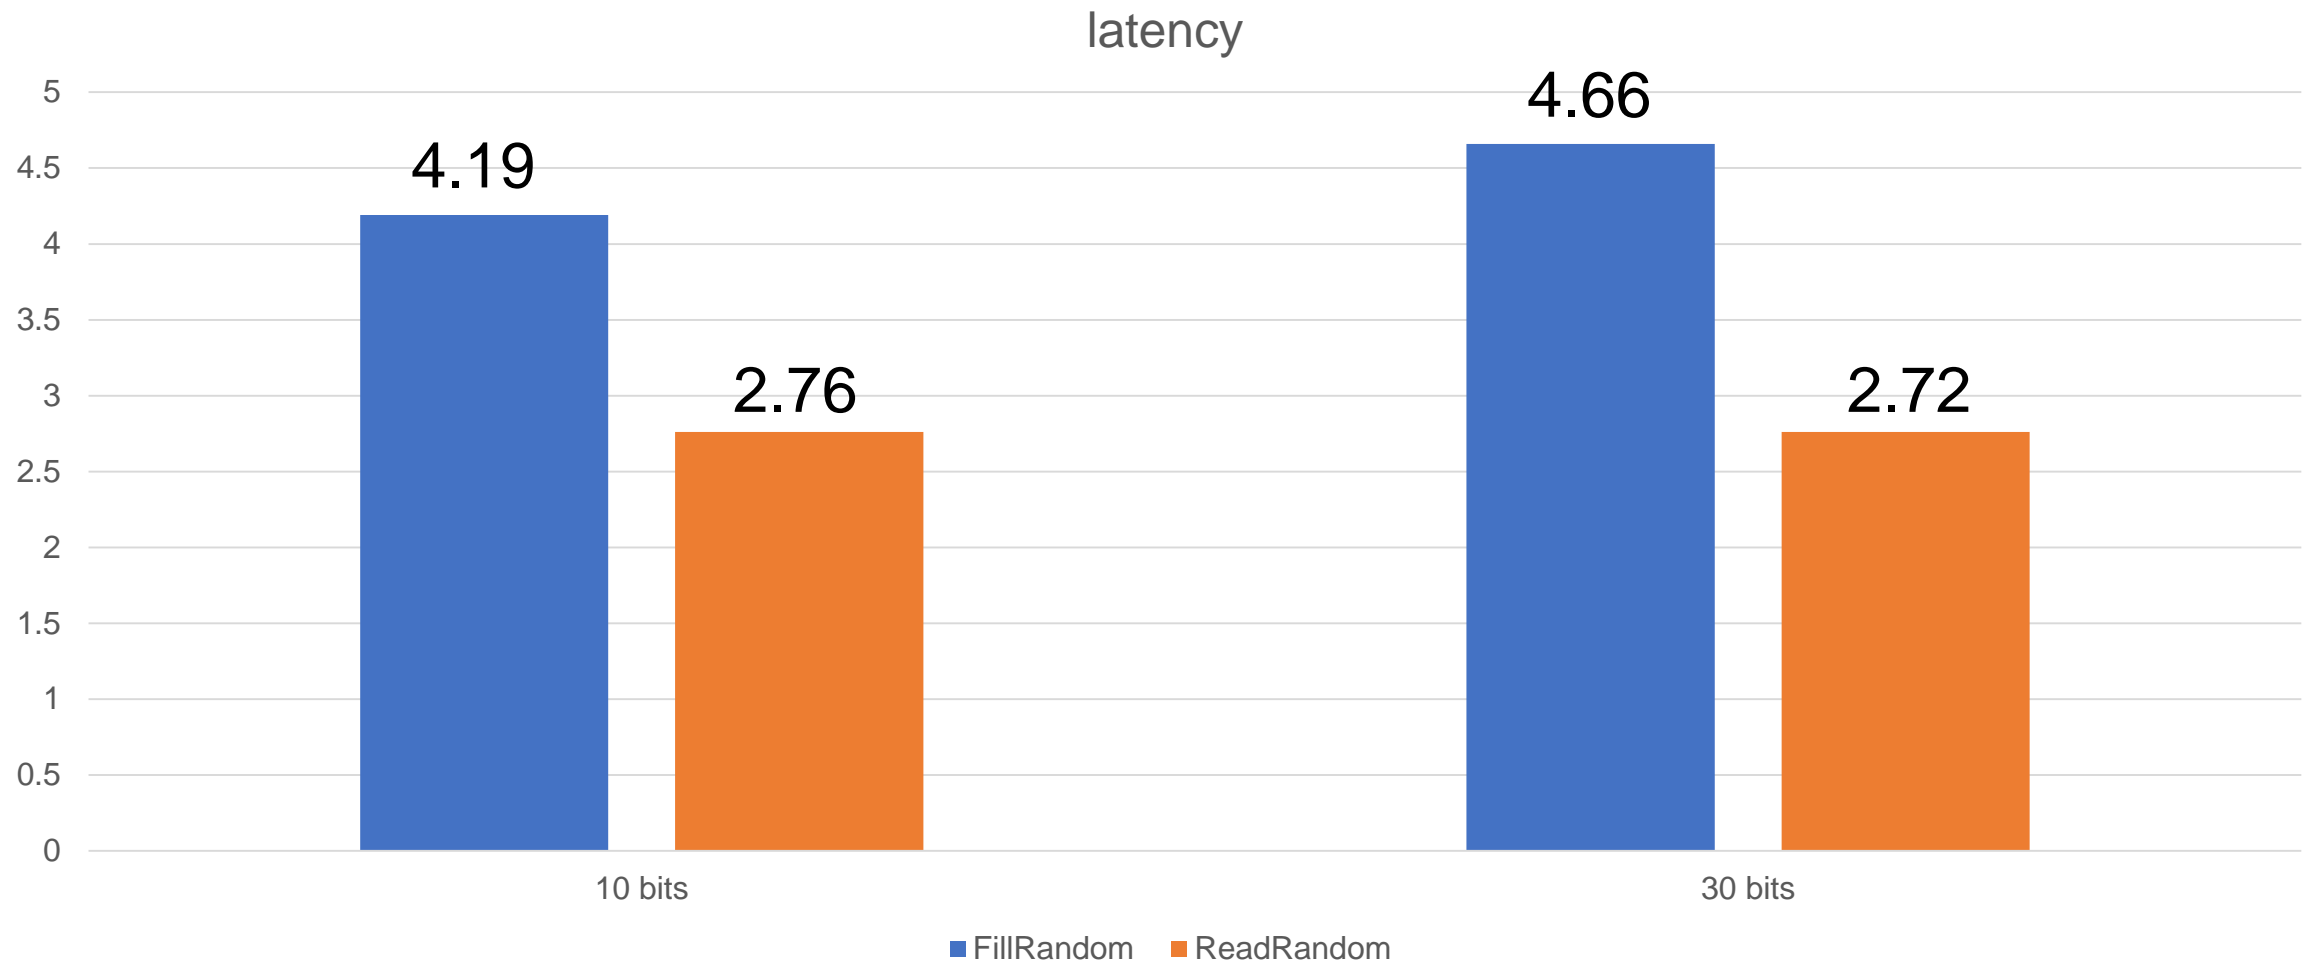
# Code edit 2

---

```
class BloomFilterPolicy : public FilterPolicy {
public:
    explicit BloomFilterPolicy(int bits_per_key) : bits_per_key_(bits_per_key) {
        // We intentionally round down to reduce probing cost a little bit
        k_ = static_cast<size_t>(bits_per_key * 0.69); // 0.69 ≈ ln(2)
        if (k_ < 1) k_ = 6;
    }
}
```



# Result



# Readmissing?

---

```
readseq      -- read N times sequentially
readreverse  -- read N times in reverse order
readrandom   -- read N times in random order
readmissing  -- read N missing keys in random order
readhot      -- read N times in random order from 1% section of DB
```

# Readmissing

```
void ReadRandom(ThreadState* thread) {
    ReadOptions options;
    std::string value;
    int found = 0;
    KeyBuffer key;
    for (int i = 0; i < reads_; i++) {
        const int k = thread->rand.Uniform(FLAGS_num);
        key.Set(k);
        if (db_->Get(options, key.slice(), &value).ok()) {
            found++;
        }
        thread->stats.FinishedSingleOp();
    }
    char msg[100];
    std::snprintf(msg, sizeof(msg), "(%d of %d found)", found,
        reads_);
    thread->stats.AddMessage(msg);
}
```

```
void ReadMissing(ThreadState* thread) {
    ReadOptions options;
    std::string value;
    KeyBuffer key;
    for (int i = 0; i < reads_; i++) {
        const int k = thread->rand.Uniform(FLAGS_num);
        key.Set(k);
        Slice s = Slice(key.slice().data(), key.slice().size() - 1);
        db_->Get(options, s, &value);
        thread->stats.FinishedSingleOp();
    }
}
```

# Readmissing

---

```
namespace {  
static uint32_t BloomHash(const Slice& key) {  
    return Hash(key.data(), key.size(), 0xbc9f1d34);  
}
```


```
uint32_t Hash(const char* data, size_t n, uint32_t seed) {  
    // Similar to murmur hash  
    const uint32_t m = 0xc6a4a793;  
    const uint32_t r = 24;  
    const char* limit = data + n;  
    uint32_t h = seed ^ (n * m);
```

# Readmissing

```
void ReadRandom(ThreadState* thread) {
    ReadOptions options;
    std::string value;
    int found = 0;
    KeyBuffer key;
    for (int i = 0; i < reads_; i++) {
        const int k = thread->rand.Uniform(FLAGS_num);
        key.Set(k);
        Slice s = Slice(key.slice().data(), key.slice().size() - 1);
        if (db_>Get(options, s, &value).ok()) {
            found++;
        }
    }
    thread->stats.FinishedSingleOp();
}
```

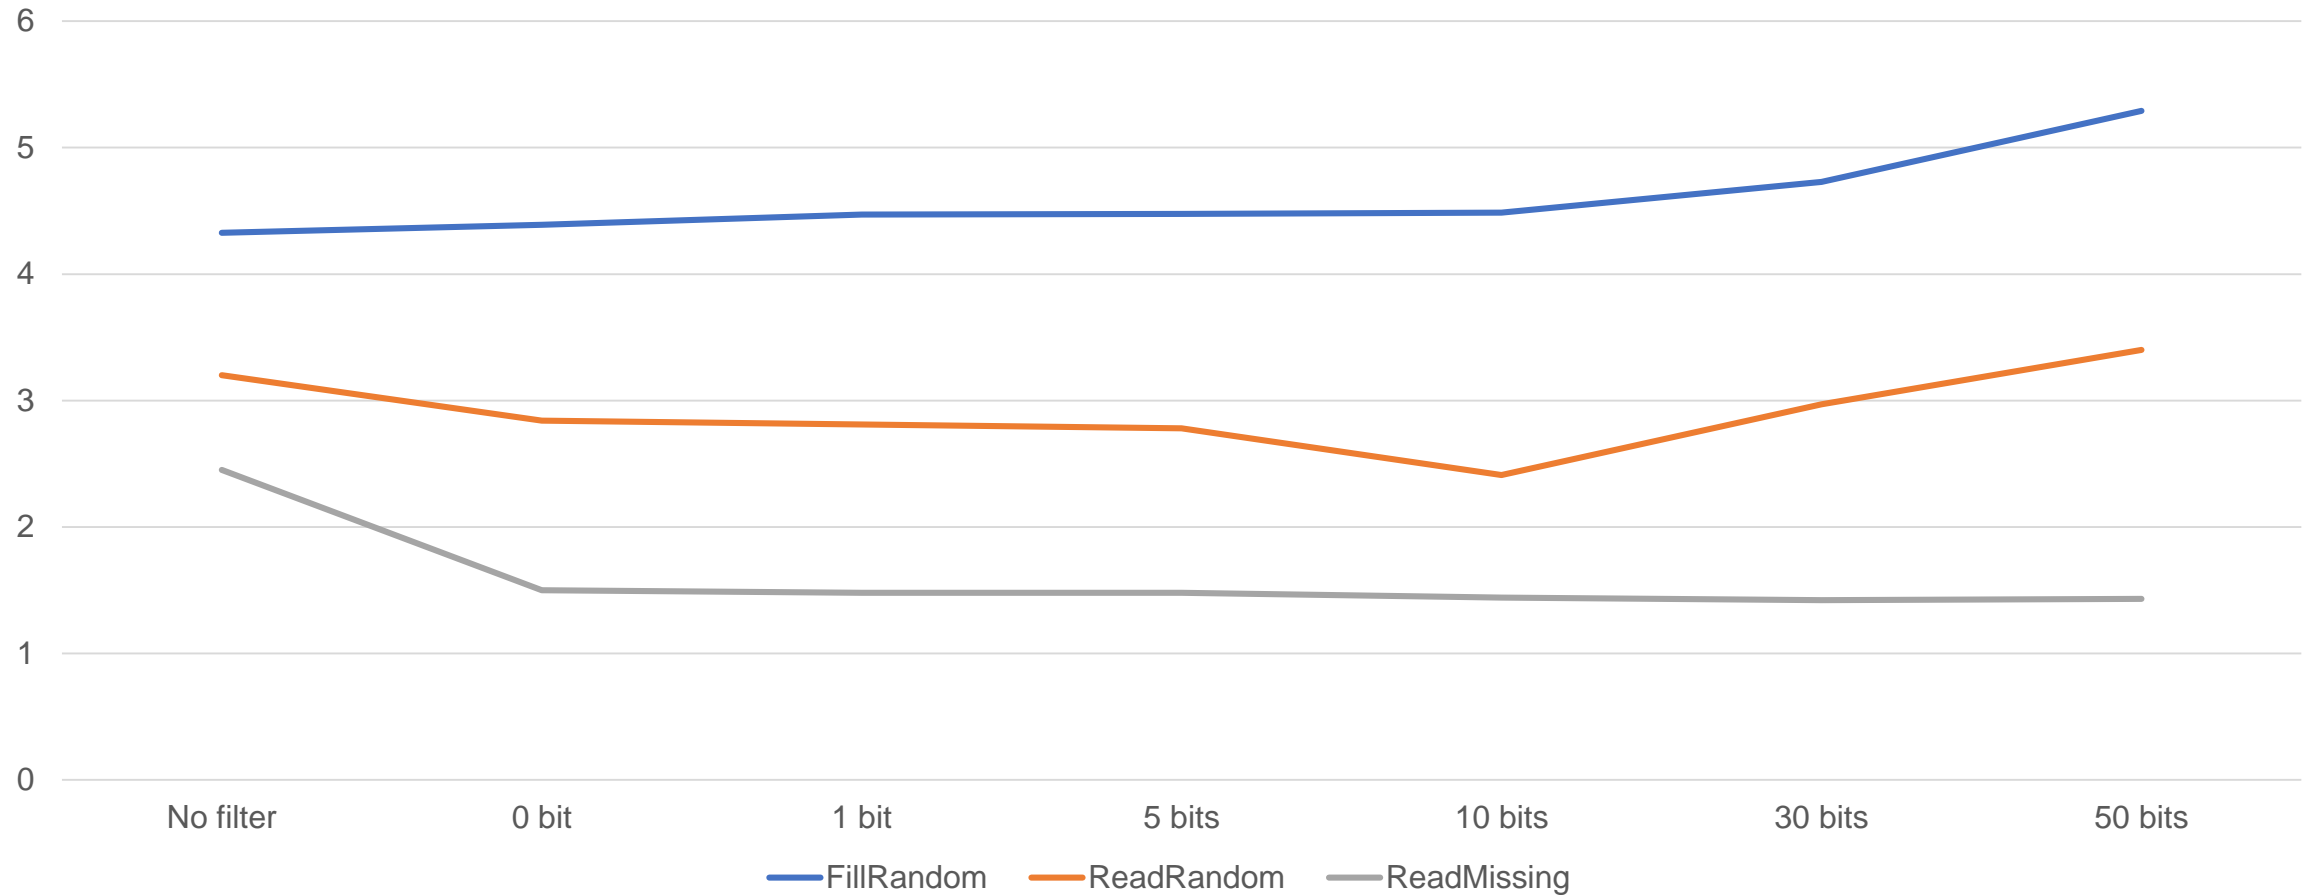
```
fillrandom    :      3.344 micros/op;   33.1 MB/s
readrandom    :      3.263 micros/op; (0 of 1000000 found)
```

# Result




# Result

Latency



# Result

Latency





# Question

---



shutterstock.com - 735394957