

# WAL / MANIFEST

Made by Isu Kim

E-Mail: [isukim16@gmail.com](mailto:isukim16@gmail.com)

# Content

- MANIFEST and Version Control
  - **CURRENT** and MANIFEST
  - VersionSet & VersionEdit
  - Initial MANIFEST generation
  - LogAndApply & EncodeTo
  - Closing DB
  - Reopening DB
- Conclusion

# CURRENT and MANIFEST

- LevelDB knows which MANIFEST file to use by CURRENT.

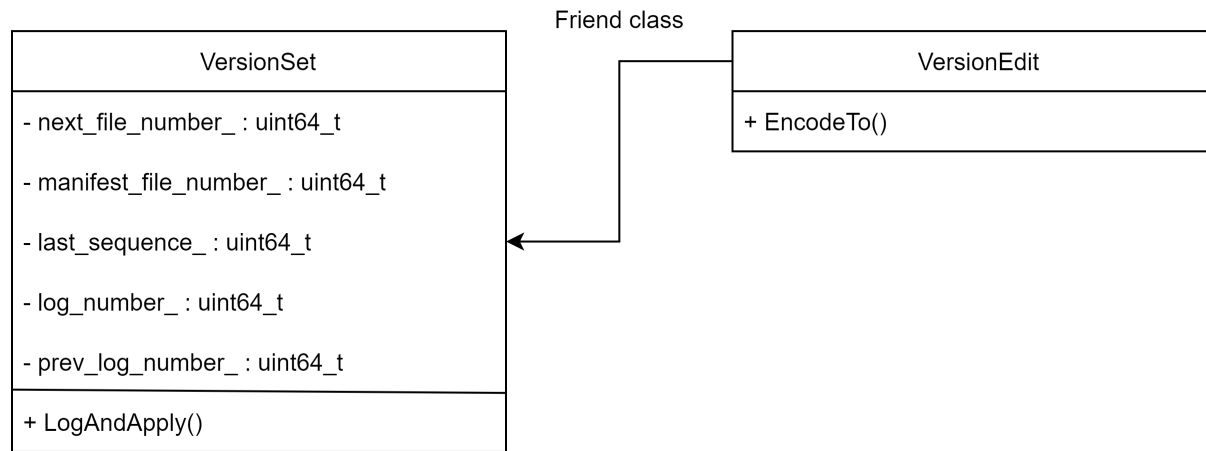
```
gooday2die@flagship:~/projects/School/2022_0.5/LevelDB/newnewnew/leveldb_debug/build$ ls -al /tmp/leveldbtest-1000/dbbench/
total 7392
drwxr-xr-x 2 gooday2die gooday2die 4096 8월 29 20:08 .
drwxr-xr-x 7 gooday2die gooday2die 4096 8월 29 20:08 ..
-rw-r--r-- 1 gooday2die gooday2die 1808353 8월 29 20:08 000005.ldb
-rw-r--r-- 1 gooday2die gooday2die 1808519 8월 29 20:08 000007.ldb
-rw-r--r-- 1 gooday2die gooday2die 2116942 8월 29 20:08 000008.log
-rw-r--r-- 1 gooday2die gooday2die 1809276 8월 29 20:08 000009.ldb
-rw-r--r-- 1 gooday2die gooday2die 16 8월 29 20:08 CURRENT
-rw-r--r-- 1 gooday2die gooday2die 0 8월 29 20:08 LOCK
-rw-r--r-- 1 gooday2die gooday2die 788 8월 29 20:08 LOG
-rw-r--r-- 1 gooday2die gooday2die 269 8월 29 20:08 MANIFEST-000002
```

- Content of CURRENT represents which MANIFEST to use.

```
gooday2die@flagship:~/projects/School/2022_0.5/LevelDB/newnewnew/leveldb_debug/build$ cat /tmp/leveldbtest-1000/dbbench/CURRENT
MANIFEST-000002
gooday2die@flagship:~/projects/School/2022_0.5/LevelDB/newnewnew/leveldb_debug/build$
```

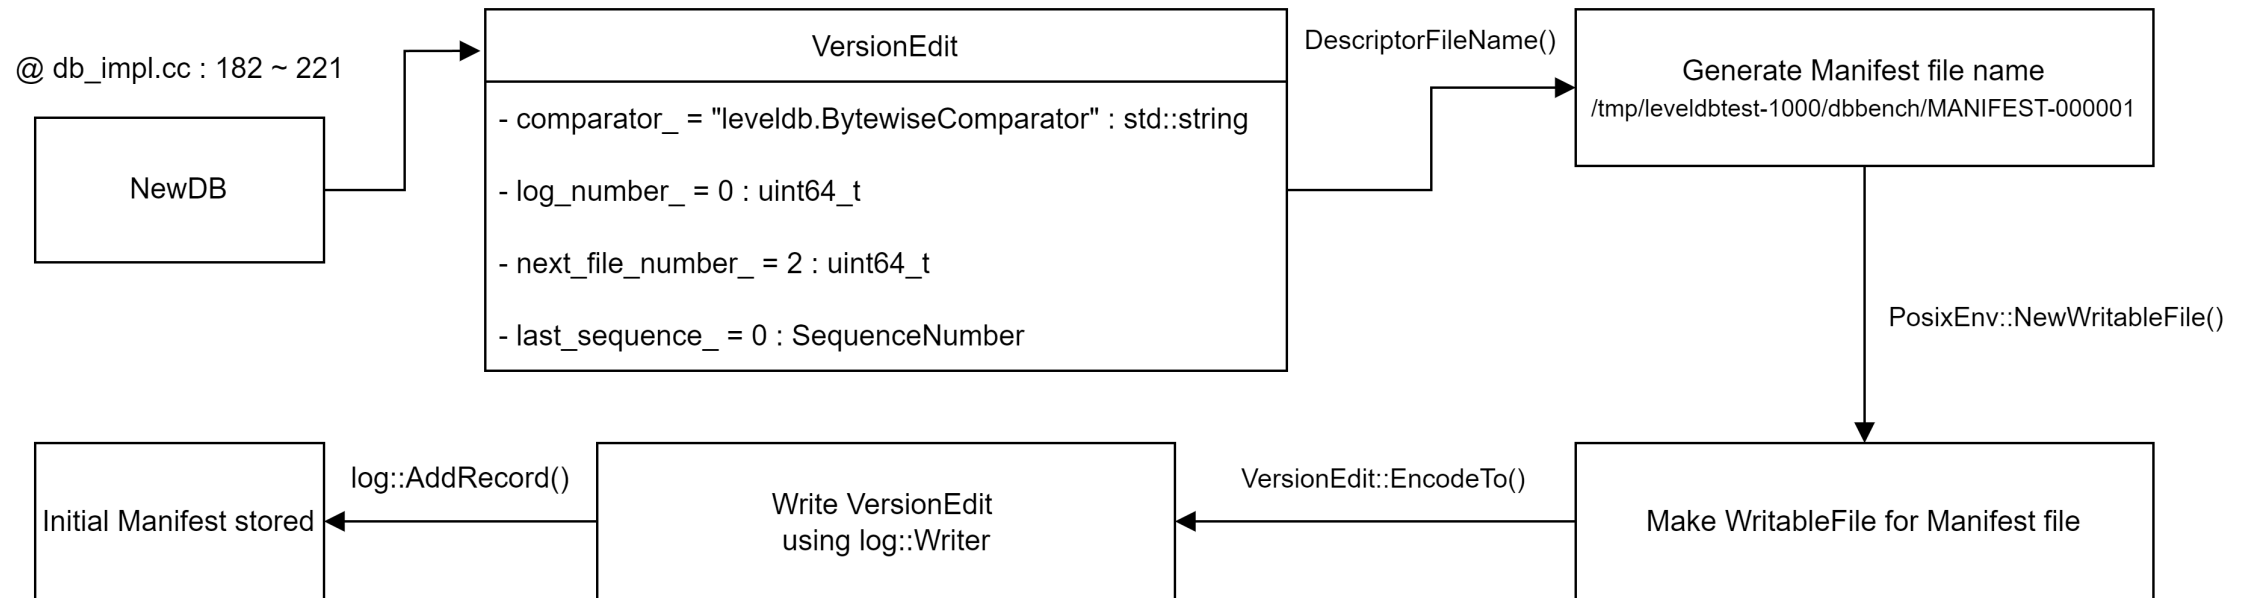
# VersionSet & VersionEdit

- VersionEdit is friend class of VersionSet
- They both store information about versions.
- LogAndApply in VersionSet, EncodeTo in VersionEdit is for recording to MANIFEST file.



# Initial MANIFEST

- When DB is being created by `DB::NewDB()`, this automatically generates MANIFEST file.
- MANIFEST file is also a `PosixEnv::WritableFile` object and uses `log::Writer`



# Initial MANIFEST - Example

- An example with initial MANIFEST
  - MANIFEST when LevelDB is being **created**. (Defined in db\_impl.cc @ 182 ~ 221)

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	0123456789ABCDEF
0000h:	56	F9	B8	F8	1C	00	01	01	1A	6C	65	76	65	6C	64	62	Vù,ø.....leveldb
0010h:	2E	42	79	74	65	77	69	73	65	43	6F	6D	70	61	72	61	.BytewiseCompara
0020h:	74	6F	72	02	00	03	02	04	00								toræ<¾.....
0030h:																	..

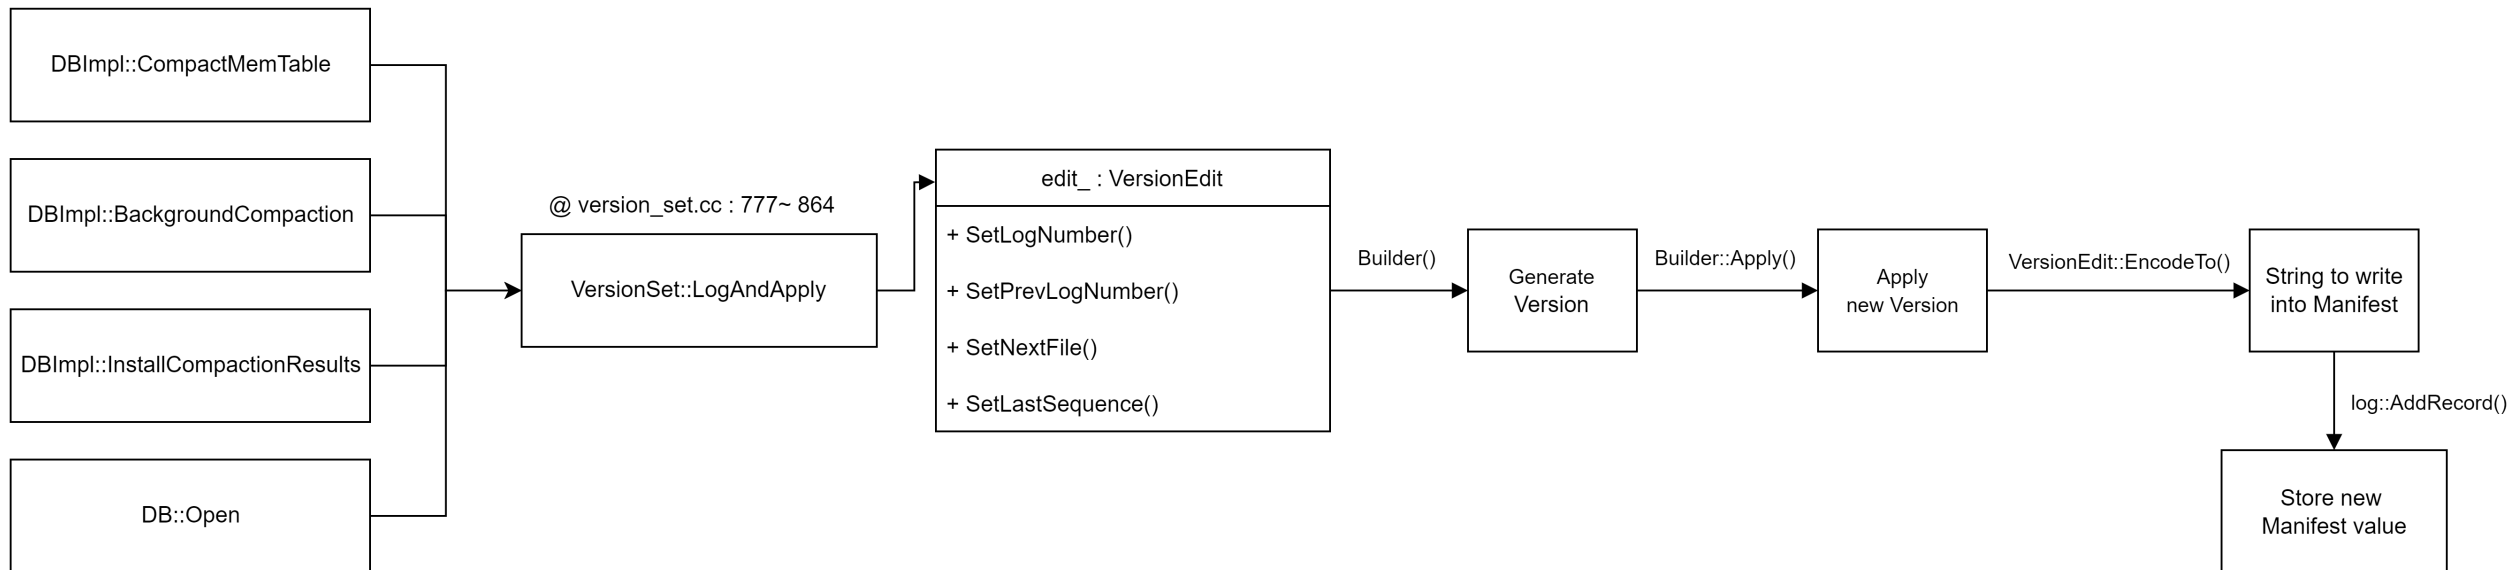
- 0x01 means this is a kComparator
- 0x1A represents total length of "leveldb.BytewiseComparator"
- This information is generated by VersionEdit::EncodeTo()

Values
Header

- |       |   |   |
|-------|---|---|
| 02 00 | ➡ | 0x02 = kLogNumber / Current log number = 0x00 (Default)     |
| 03 02 | ➡ | 0x03 = kNextFileNumber / Next file number is 0x02 (Default) |
| 04 00 | ➡ | 0x04 is kLastSequence / Last sequence was 0x00 (Default)    |

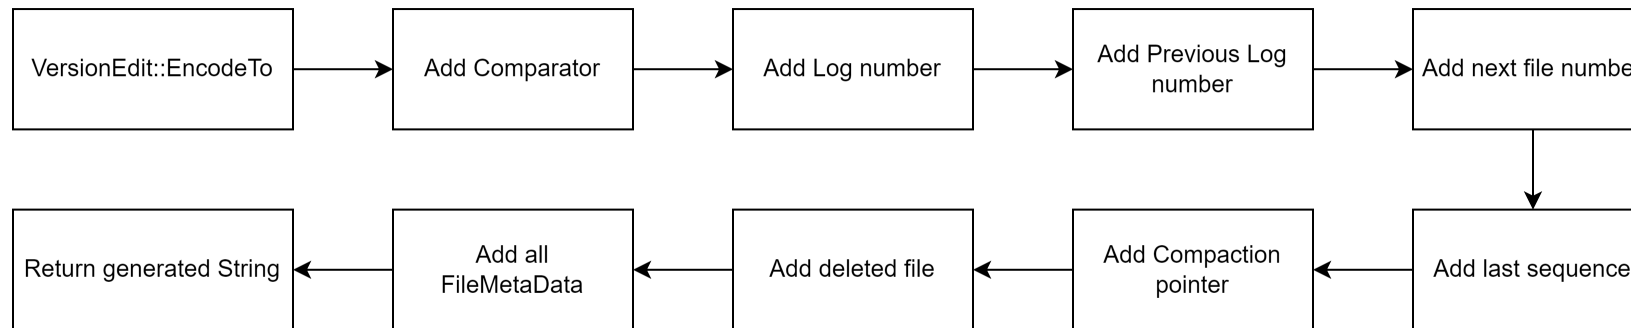
# VersionSet::LogAndApply

- LogAndApply is for version control.
- Is called in following member functions.
- Stores data to MANIFEST file.
- DBImpl::CompactMemtable, DBImpl::BackgroundCompaction, DBImpl::InstallCompactionResults, DB::Open calls this



# VersionEdit::EncodeTo

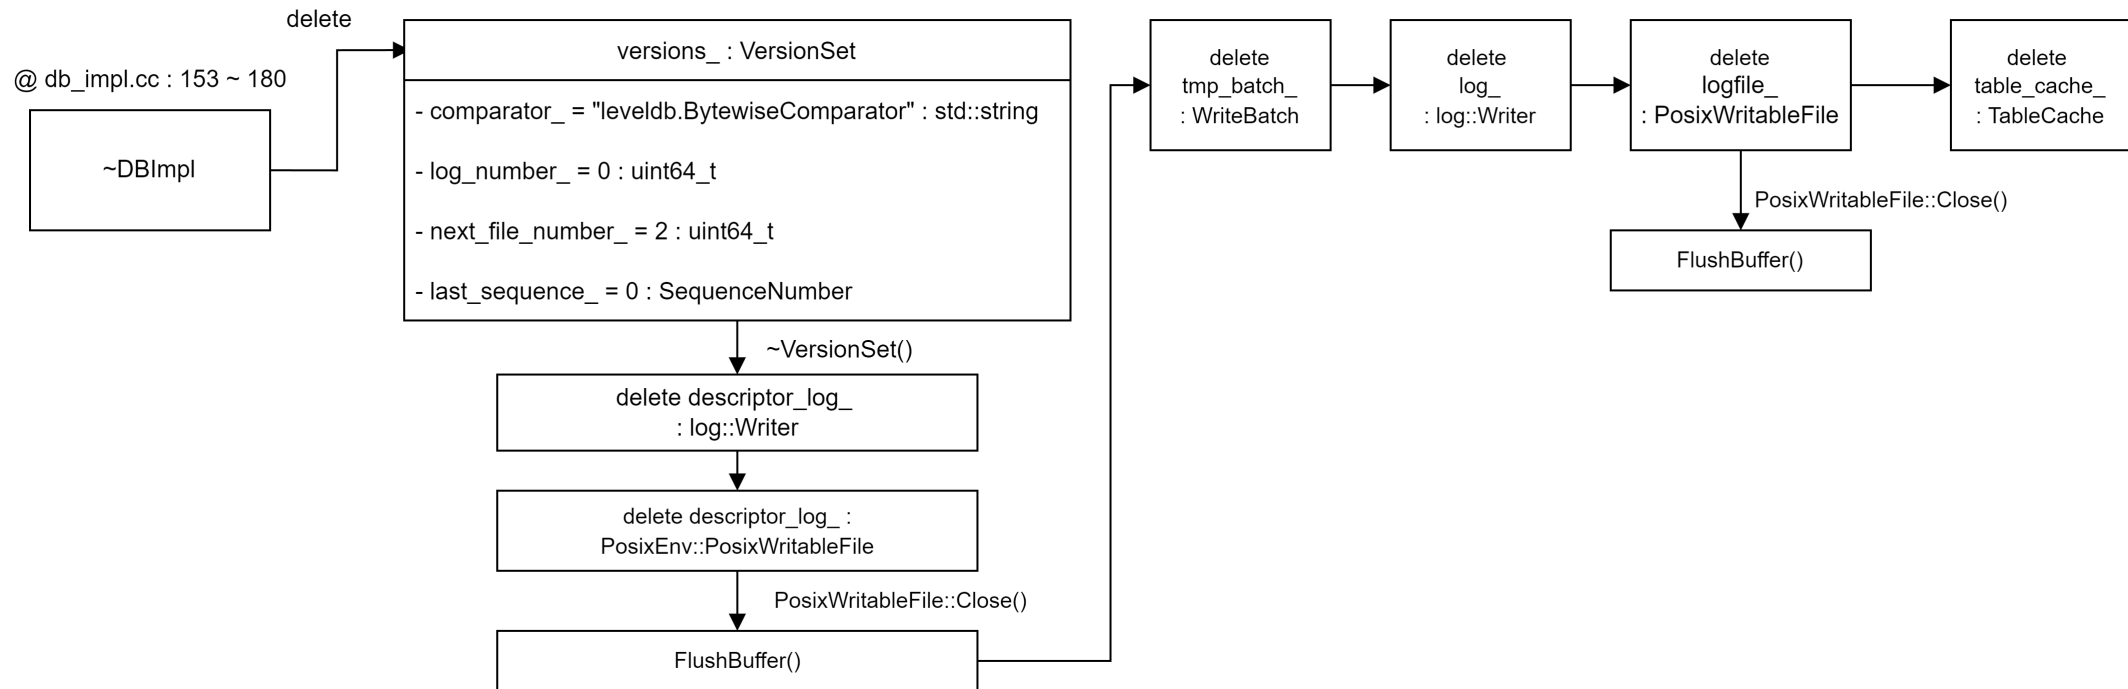
- Generates actual data to write into MANIFEST file from VersionSet and FileMetaData
- By the generated string, the MANIFEST file is written.

[illegible]



# Closing DB

- ~DBImpl is called when closing DB (delete)



# Closing DB

- Let's see how manifest file is written with fillseq, num == 10
- Similar to previous file.



```
0020h: 74 6F 72 A4 9C 8B BE 08 00 01 02 03 09 00 03 04 toræ¼.....
0030h: 04 00 ..
```

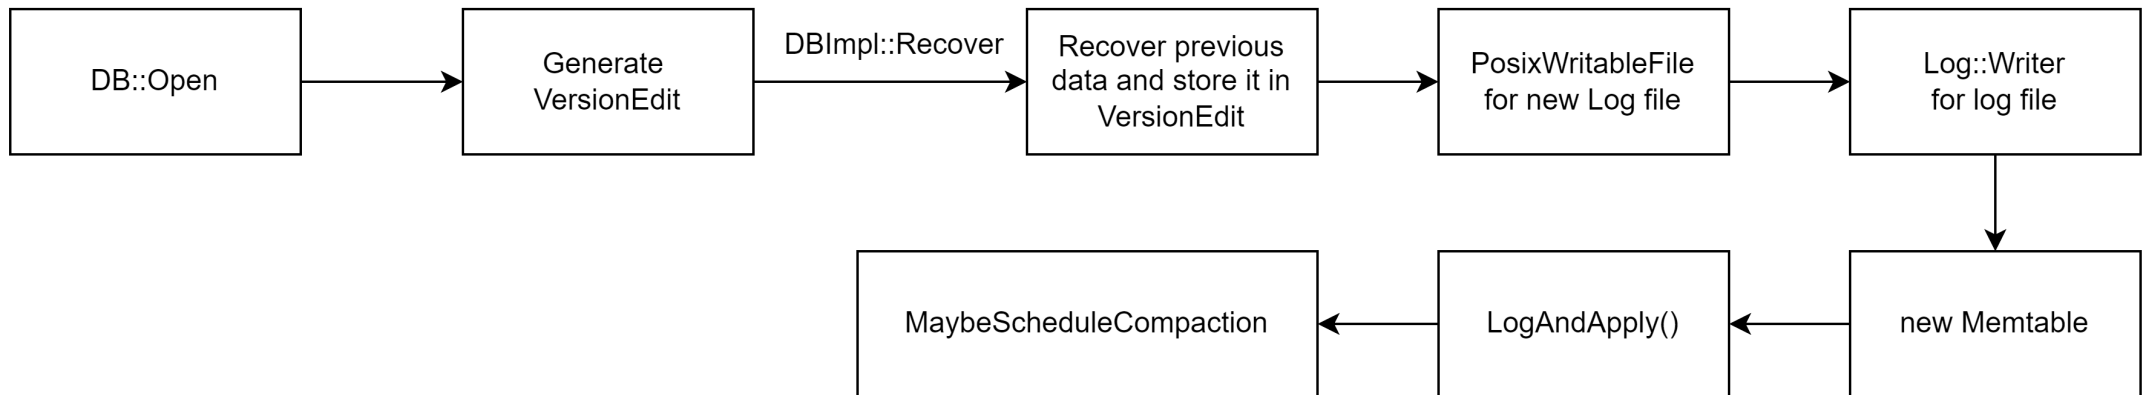
- |              |   |
|--------------|---|
| <b>02 03</b> | ➡ 0x02 is kLogNumber & Current log number is 0x03 (Current .log file's number is 000003)                          |
| <b>09 00</b> | ➡ 0x09 is kPrevLogNumber & Previous log number is 0x00  |
| <b>03 04</b> | ➡ 0x03 is kNextFileNumber & Next file number is 0x04 (When next MANIFEST file is generated, the number will be 4) |
| <b>04 00</b> | ➡ 0x04 is kLastSequence & Last sequence was 0x00 (This will be updated in next LogAndApply)                       |

```
[+] LogNumber : 3
[+] Previous Log number : 0
[+] Next file : 4
[+] Last Sequence : 0
[+] EncodeTo was called!
[+] DST VALUE (LogNumber): 02 03
[+] DST VALUE (PrevLogNumber): 02 03 09 00
[+] DST VALUE (NextFileNumber): 02 03 09 00 03 04
[+] DST VALUE (LastSequence): 02 03 09 00 03 04 04 00
[+] Record value: new record to MANIFEST: 02 03 09 00 03 04 04 00
```

# Opening DB

- Let's see how manifest file is written with fillseq, num == 10, use\_existing\_db=1
- Opening DB calls LogAndApply!
- Thus, MANIFEST will be changed.

@ db\_impl.c : 1510 ~ 1551



# Opening DB - Manifest

- MANIFEST changed
- CURRENT was set to MANIFEST-000004

MANIFEST-000002 x

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	0123456789ABCDEF
0000h:	56	F9	B8	F8	1C	00	01	01	1A	6C	65	76	65	6C	64	62	Vù,ø.....leveldb
0010h:	2E	42	79	74	65	77	69	73	65	43	6F	6D	70	61	72	61	.BytewiseCompara
0020h:	74	6F	72	A4	9C	8B	BE	08	00	01	02	03	09	00	03	04	toræ¼.....
0030h:	04	00															..

MANIFEST-000004 x

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	0123456789ABCDEF
0000h:	56	F9	B8	F8	1C	00	01	01	1A	6C	65	76	65	6C	64	62	Vù,ø.....leveldb
0010h:	2E	42	79	74	65	77	69	73	65	43	6F	6D	70	61	72	61	.BytewiseCompara
0020h:	74	6F	72	CE	6F	F5	52	3F	00	01	02	06	09	00	03	07	torÎøR?.....
0030h:	04	0A	07	00	05	EF	05	18	30	30	30	30	30	30	30	30	.....ï..00000000
0040h:	30	30	30	30	30	30	30	30	01	01	00	00	00	00	00	00	00000000.....
0050h:	18	30	30	30	30	30	30	30	30	30	30	30	30	30	30	30	.0000000000000000
0060h:	39	01	0A	00	00	00	00	00	00								9.....

Before

After

- |       |   |   |
|-------|---|---|
| 02 06 | ➡ | 0x02 is kLogNumber & Current log number is 0x06 (00006.log is the WAL file)                                     |
| 09 00 | ➡ | 0x09 is kPrevLogNumber & Previous log number is 0x00  |
| 03 04 | ➡ | 0x03 is kNextFileNumber & Next file number is 0x07 (When next MANIFEST file is generated, the number will be 7) |
| 04 00 | ➡ | 0x04 is kLastSequence & Last sequence was 0x0A (Yes, we had 10 entries that were put)                           |

Values

Header

# Opening DB - Manifest

- MANIFEST changed
- CURRENT was set to MANIFEST-000004

MANIFEST-000002

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	0123456789ABCDEF
0000h:	56	F9	B8	F8	1C	00	01	01	1A	6C	65	76	65	6C	64	62	Vù,ø.....leveldb
0010h:	2E	42	79	74	65	77	69	73	65	43	6F	6D	70	61	72	61	.BytewiseCompara
0020h:	74	6F	72	A4	9C	8B	BE	08	00	01	02	03	09	00	03	04	toræ¾.....
0030h:	04	00															..

MANIFEST-000004

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	0123456789ABCDEF
0000h:	56	F9	B8	F8	1C	00	01	01	1A	6C	65	76	65	6C	64	62	Vù,ø.....leveldb
0010h:	2E	42	79	74	65	77	69	73	65	43	6F	6D	70	61	72	61	.BytewiseCompara
0020h:	74	6F	72	CE	6F	F5	52	3F	00	01	02	06	09	00	03	07	torIoøR?.....
0030h:	04	0A	07	00	05	EF	05	18	30	30	30	30	30	30	30	30	.....i..00000000
0040h:	30	30	30	30	30	30	30	30	01	01	00	00	00	00	00	00	00000000.....
0050h:	18	30	30	30	30	30	30	30	30	30	30	30	30	30	30	30	.0000000000000000
0060h:	39	01	0A	00	00	00	00	00	00								9.....

07

➡ 0x07 is kNewFile

00

➡ 0x00 is the current level of LevelDB (We have total 10 entries of 116 bytes, which resides in level 0)

05

➡ 0x05 is FileMetaData's number

04 00

➡ 0x04 is kLastSequence & Last sequence was 0x0A (Yes, we had 10 entries that were put)

EF 05

➡ 0xef05 is the size of the file.

18 ~ 00

➡ The smallest FileMetaData's key value encoded. (Min key value was 0000000000000000) since this was fillseq

18 ~ 00

➡ The largest FileMetaData's key value encoded. (Max key value was 0000000000000009) since this was fillseq

Header

# Conclusion

---

- LogAndApply, EncodeTo takes care of version control.
- These functions generate MANIFEST file.
- Research how following functions work with MANIFEST file.
  - CompactMemTable
  - BackgroundCompaction
  - InstallCompactionResults

# Final Conclusion

---

- Regrets on too much focus on WAL not Version & MANIFEST.
- Couldn't see the forest for the trees

