

Team Compaction

강상우

E-Mail : aarom416@naver.com

발표: 좌우꾸와썬


E-Mail : erosbryant@dankook.ac.kr

Contents

- TwoLevelIterator
- MergingIterator

TwoLevelIterator

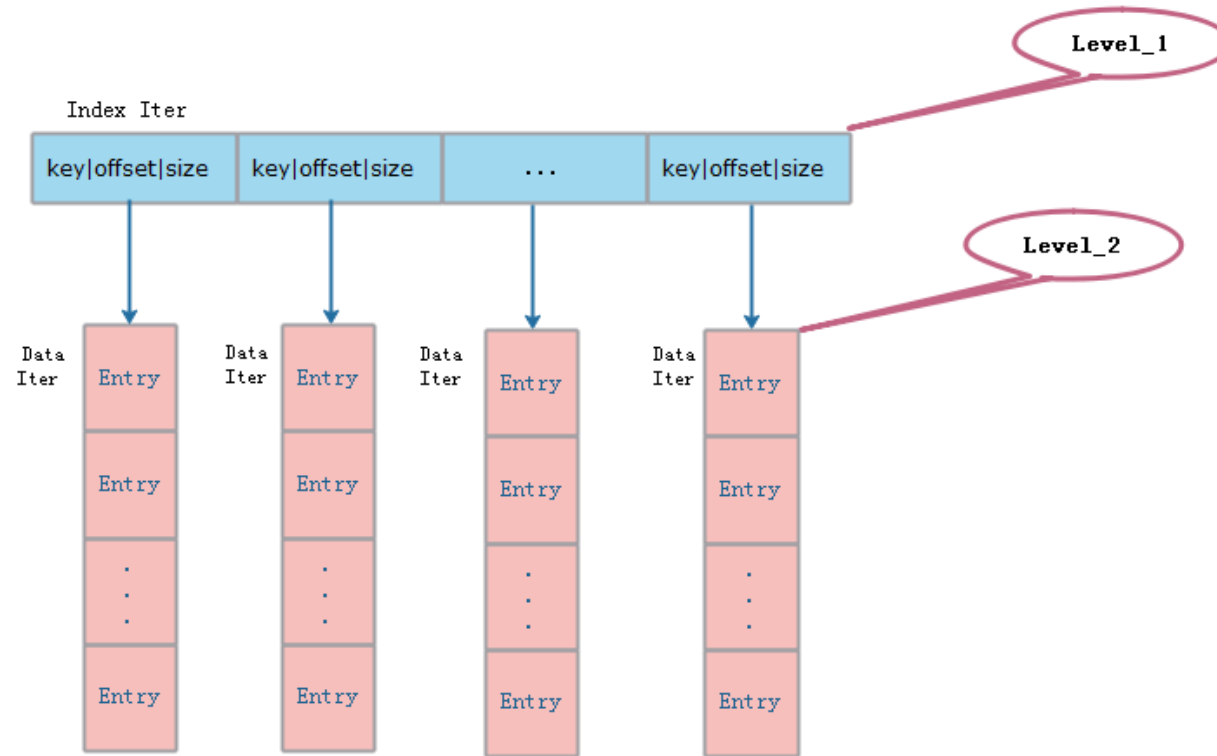
- Using in the SST



<https://leveldb-handbook.readthedocs.io/zh/latest/sstable.html>

<https://leveldb-handbook.readthedocs.io/zh/latest/sstable.html>

TwoLevelIterator



https://img-blog.csdnimg.cn/20200614175126109.png#pic_center?x-oss-process=image/watermark,type_ZmFuZ3poZW5naGVpdGk,shadow_10,text_aHR0cHM6Ly9ibG9nLmNzZG4ubmV0L0g1MTQ0MzQ0ODU=,size_16,color_FFFFFFFF,t_70

TwoLevelIterator

Iterator class
상속

```
class TwoLevelIterator : public Iterator {  
public:  
    TwoLevelIterator(Iterator* index_iter, BlockFunction block_function,  
                    void* arg, const ReadOptions& options);  
};
```

SST의 경우:

1. index_iter는 index block을 가리키는 iter
2. block_function은 Table::BlockReader, 즉 블록을 읽는 것
3. arg는 SST을 가리키는 것
4. Options 읽기 옵션

MergingIterator

- MergingIterator
 - 이름 대로 Iterator들을 merge
- NewInternalIterator
 - 핵심 함수

```
Iterator* DBImpl::NewInternalIterator(const ReadOptions& options,  
                                     SequenceNumber* latest_snapshot,  
                                     uint32_t* seed) {  
    mutex_.Lock();  
    *latest_snapshot = versions_->LastSequence();  
  
    // Collect together all needed child iterators  
    std::vector<Iterator*> list;  
    list.push_back(mem_->NewIterator());  
    mem_->Ref();  
    if (imm_ != nullptr) {  
        list.push_back(imm_->NewIterator());  
        imm_->Ref();  
    }  
    versions_->current()->AddIterators(options, &list);  
    Iterator* internal_iter =  
        NewMergingIterator(&internal_comparator_, &list[0], list.size());  
    versions_->current()->Ref();  
  
    IterState* cleanup = new IterState(&mutex_, mem_, imm_, versions_->current());  
    internal_iter->RegisterCleanup(CleanupIteratorState, cleanup, nullptr);  
  
    *seed = ++seed_;  
    mutex_.Unlock();  
    return internal_iter;  
}
```

MergingIterator

Vector<iterator*>



Memtable

Imm Memtable

level 0

level 1 ~ N

```
void Version::AddIterators(const ReadOptions& options,  
    std::vector<Iterator*>* iters) {  
    // Merge all level zero files together since they may overlap  
    for (size_t i = 0; i < files_[0].size(); i++) {  
        iters->push_back(vset->table cache->NewIterator(  
            options, files_[0][i]->number, files_[0][i]->file_size));  
    }  
  
    // For levels > 0, we can use a concatenating iterator that sequentially  
    // walks through the non-overlapping files in the level, opening them  
    // lazily.  
    for (int level = 1; level < config::kNumLevels; level++) {  
        if (!files_[level].empty()) {  
            iters->push_back(NewConcatenatingIterator(options, level));  
        }  
    }  
}
```

MergingIterator

Vector<iterator*>



Memtable

Imm Memtable

level 0

level 1 ~ N

```
void Version::AddIterators(const ReadOptions& options,  
    std::vector<Iterator*>* iters) {  
    // Merge all level zero files together since they may overlap  
    for (size_t i = 0; i < files_[0].size(); i++) {  
        iters->push_back(vset->table cache->NewIterator(  
            options, files_[0][i]->number, files_[0][i]->file_size));  
    }  
  
    // For levels > 0, we can use a concatenating iterator that sequentially  
    // walks through the non-overlapping files in the level, opening them  
    // lazily.  
    for (int level = 1; level < config::kNumLevels; level++) {  
        if (!files_[level].empty()) {  
            iters->push_back(NewConcatenatingIterator(options, level));  
        }  
    }  
}
```


MergingIterator

- MergingIterator
 - 이름 대로 Iterator들을 merge
- NewInternalIterator
 - 핵심 함수

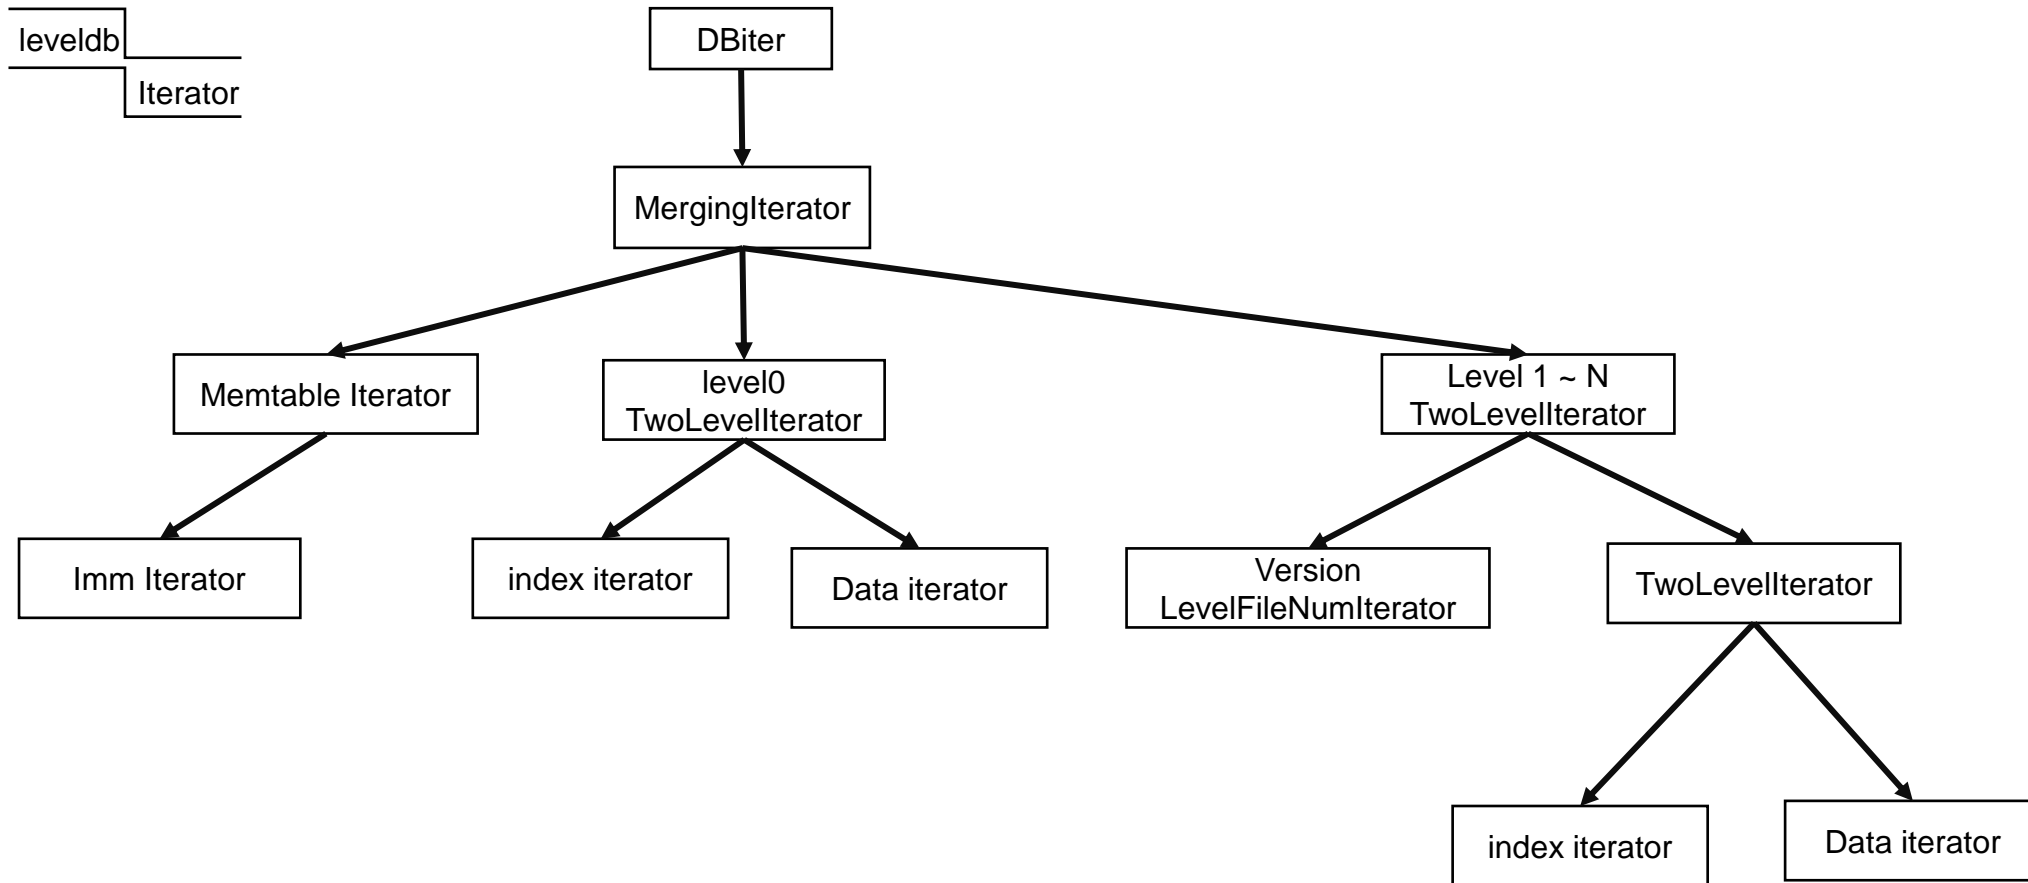
```
Iterator* DBImpl::NewInternalIterator(const ReadOptions& options,
                                     SequenceNumber* latest_snapshot,
                                     uint32_t* seed) {
    mutex_.Lock();
    *latest_snapshot = versions_->LastSequence();

    // Collect together all needed child iterators
    std::vector<Iterator*> list;
    list.push_back(mem_->NewIterator());
    mem_->Ref();
    if (imm_ != nullptr) {
        list.push_back(imm_->NewIterator());
        imm_->Ref();
    }
    versions_->current()->AddIterators(options, &list);
    Iterator* internal_iter =
        NewMergingIterator(&internal_comparator_, &list[0], list.size());
    versions_->current()->Ref();

    IterState* cleanup = new IterState(&mutex_, mem_, imm_, versions_->current());
    internal_iter->RegisterCleanup(CleanupIteratorState, cleanup, nullptr);

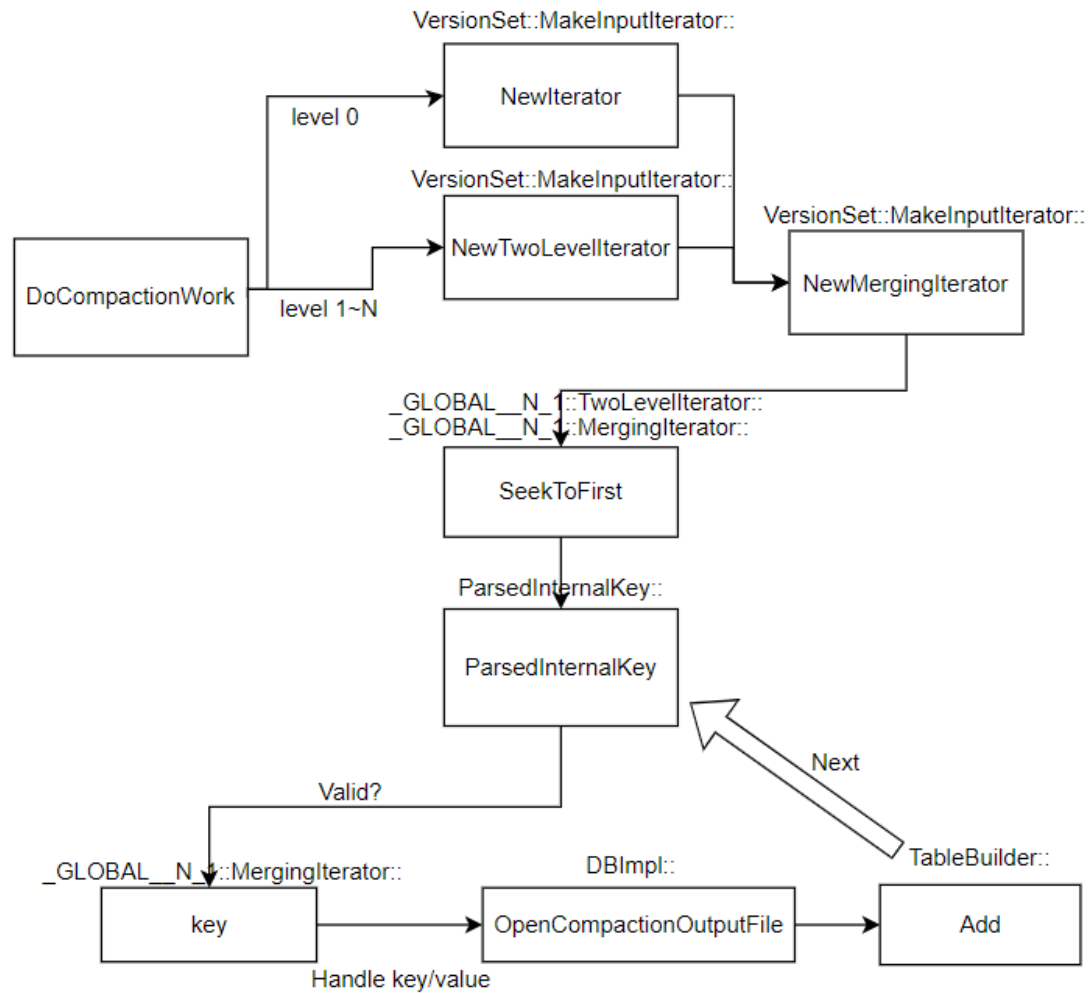
    *seed = ++seed_;
    mutex_.Unlock();
    return internal_iter;
}
```

Iterator 기능적 흐름도



<https://www.jianshu.com/p/9e63dea36af0>

DoCompactionWork



Thank you