

# Week3 Homework Practice 1 & 3

2022. 7. 19

Suhwan Sin

E-Mail: [tlstnghks77@dankook.ac.kr](mailto:tlstnghks77@dankook.ac.kr)

# Content

- Practice 1
  - Meta operations
  - Option (fillseq, fillrandom)
  - Compare stats
  - Q1
  - Q2
  - Q3
  
- Practice 3
  - Batch processing
  - Conclusion

# Practice 1

---

[A] \$ ./db\_bench --benchmarks="fillseq"

[B] \$ ./db\_bench --benchmarks="fillrandom"

Q1. Compare throughput, latency, and stats of two benchmarks and explain why.

Q2. In benchmark A, SSTs are not written in L0. Why?

Q3. Calculate SAF (Space Amplification Factor) for each benchmark.

# Meta operations

# compact

`compact`      -- Compact the entire DB

```
sh@ssh-desktop:~/leveldb/build$ ./db_bench --benchmarks="fillrandom,compact"
compact      : 802704.000 micros/op;
```

# stats

`stats`        -- Print DB stats

```
sh@ssh-desktop:~/leveldb/build$ ./db_bench --benchmarks="fillrandom,stats"
```

Compactions					
Level	Files	Size(MB)	Time(sec)	Read(MB)	Write(MB)
0	5	16	0	0	100
1	11	20	1	125	115
2	25	48	1	119	114

# Meta operations

# sstables

sstables -- Print sstable info

```
sh@ssh-desktop:~/leveldb/build$ ./db_bench --benchmarks="fillrandom,sstables"
```

--- level 0 ---

186:3272231['00000000000000010' @ 937527 : 1 .. '00000000000999929' @ 937791 : 1]

--- level 1 ---

183:2117005['00000000000080482' @ 775408 : 1 .. '00000000000124952' @ 555648 : 1]

184:2116987['00000000000124953' @ 914465 : 1 .. '00000000000169090' @ 825729 : 1]

187:2117005['00000000000169092' @ 503059 : 1 .. '00000000000213145' @ 829366 : 1]

--- level 2 ---

201:2119733['00000000000000000' @ 854569 : 1 .. '0000000000030480' @ 634990 : 1]

202:2115634['0000000000030481' @ 295655 : 1 .. '0000000000060833' @ 4325 : 1]

203:2116514['0000000000060836' @ 410002 : 1 .. '0000000000099982' @ 10709 : 1]

# Option – “fillseq & fillrandom”

```
fillseq      -- write N values in sequential key order in async mode
fillrandom   -- write N values in random key order in async mode
```

	Key range 중복	Compaction	Example		
fillseq	x	x	1 2 3	4 5 6	7 8 9
fillrandom	o	o	3 7 2	5 4 3	1 1 1

# Compare Seq / Random

fillseq : 1.190 micros/op; 93.0 MB/s

Compactions

Level	Files	Size(MB)	Time(sec)	Read(MB)	Write(MB)
<hr/>					
2	33	102	0	0	108
3	2	6	0	0	0

fillrandom : 1.889 micros/op; 58.6 MB/s

Compactions

Level	Files	Size(MB)	Time(sec)	Read(MB)	Write(MB)
<hr/>					
0	6	19	0	0	103
1	12	22	1	125	115
2	24	47	0	114	109

# Q1.

---

**Q1. Compare throughput, latency, and stats of two benchmarks and explain why.**

**A.**

	throughput	latency
fillseq	↑	↑
fillrandom	↓	↓



# Q2.

---

**Q2. In benchmark A, SSTs are not written in L0. Why?**

**A. Trivial move**

LevelDB One kind of optimization is when the following conditions are met

- There is only one file in the layer
- Layer files and level+1, Layer files do not overlap
- Layer files and level+2, The file size of the file overlapping part of the layer does not exceed the threshold


**The file of the layer is moved to level+1 Delamination **directly****

# Q2. Trivial move


---

- **Fillseq**
  - L2, L2, L2, L2 ... (Compaction) L3
- **Fillrandom**
  - L2, L1, L0, L0 (Compaction) ... L0

## Q2. Trivial move



## Q2. Trivial move



## Q2. Trivial move

4, 5, 6    7, 8, 9    10, 11, 12

MEM

L0

L1

L2

1, 2, 3

fillseq

2, 5, 8    4, 6, 9    10, 11, 12

MEM

L0


L1

L2


7, 3, 1

fillrandom


## Q2. Trivial move



## Q2. Trivial move



## Q2. Trivial move





# Q3.

## Q3. Calculate SAF (Space Amplification Factor) for each benchmark.

$$A. (102 + 6) / 110 = 0.98$$

```
./db_bench --benchmarks="fillseq,stats,compact,stats"
```

```
fillseq      :      1.174 micros/op;   94.2 MB/s
```

Level	Files	Size(MB)	Compactions		
			Time(sec)	Read(MB)	Write(MB)
2	33	102	0	0	108
3	2	6	0	0	0

```
compact      : 727065.000 micros/op;
```

Level	Files	Size(MB)	Compactions		
			Time(sec)	Read(MB)	Write(MB)
2	0	0	0	0	110
3	67	110	1	97	97

$$A. (19 + 22 + 47) / 69 = 1.275$$

```
./db_bench --benchmarks="fillrandom,stats,compact,stats"
```

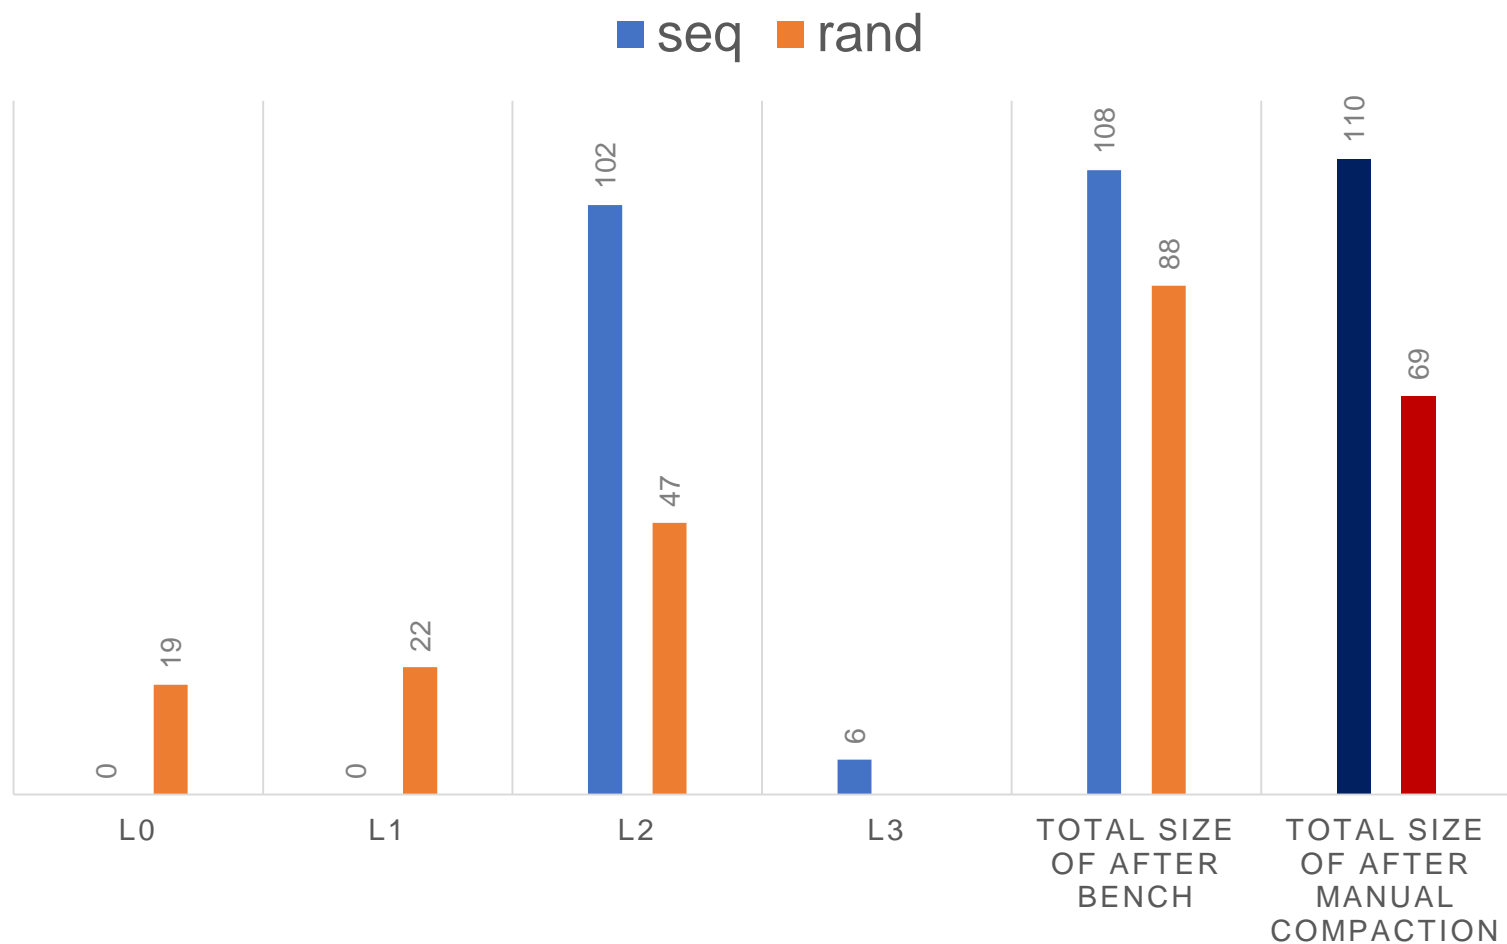
```
fillrandom   :      1.948 micros/op;   56.8 MB/s
```

Level	Files	Size(MB)	Compactions		
			Time(sec)	Read(MB)	Write(MB)
0	6	19	0	0	103
1	12	22	1	125	115
2	24	47	0	114	109

```
compact      : 923866.000 micros/op;
```

Level	Files	Size(MB)	Compactions		
			Time(sec)	Read(MB)	Write(MB)
0	0	0	0	0	104
1	0	0	1	166	150
2	36	69	1	230	210

# Q3.



	size	
level	Seq (MB)	Rand (MB)
L0	0	19
L1	0	22
<u>L2</u>	102	47
L3	6	
Total size of After Bench	108	88
Total size of After Manual Compaction	110	69
<b>SAF</b>	<b>0.98</b>	<b>1.275</b>

# Practice 3

---

[A] \$ ./db\_bench --benchmarks="fillrandom" --value\_size=100 --num=1000000 --compression\_ratio=1

[B] \$ ./db\_bench --benchmarks="fillrandom" --value\_size=1000 --num=114173 --compression\_ratio=1

Note 1. key\_size = 16 Bytes

Note 2. same total kv pairs size.


Note 3. # of B's entries =  $114173 = (16+100)/(16+1000) * 1000000$

Q. The size of input kv pairs is the same. But one has better throughput and latency than the other.

Explain why.

# Batch processing

Batch Processing is a way to process the job in groups



	DB size	# of entries	Entry size	Throughput	Latency
A	Same	1,000,000	116B	↓	↑
B	Same	114173	1016B	↑	↓

# Appendix

## - Trivial move in “fillseq”

```
sh@ssh-desktop:~/leveldb/build$ ./db_bench --benchmarks="fillseq,stats" --num=70000
LevelDB: version 1.23
Date: Tue Jul 19 09:30:16 2022
CPU: 16 * Intel(R) Core(TM) i7-10700K CPU @ 3.80GHz
CPUTCache: 16384 KB
Keys: 16 bytes each
Values: 100 bytes each (50 bytes after compression)
Entries: 70000
RawSize: 7.7 MB (estimated)
FileSize: 4.4 MB (estimated)
WARNING: Snappy compression is not enabled
```

```
-----
fillseq      :      1.330 micros/op;   83.2 MB/s
```

	Level	Files	Size(MB)	Time(sec)	Read(MB)	Write(MB)
--	-------	-------	----------	-----------	----------	-----------

	2	2	6	0	0	6
--	---	---	---	---	---	---

```
sh@ssh-desktop:~/leveldb/build$ ./db_bench --benchmarks="fillseq,stats" --num=1000000
LevelDB: version 1.23
Date: Tue Jul 19 09:30:35 2022
CPU: 16 * Intel(R) Core(TM) i7-10700K CPU @ 3.80GHz
CPUTCache: 16384 KB
Keys: 16 bytes each
Values: 100 bytes each (50 bytes after compression)
Entries: 1000000
RawSize: 110.6 MB (estimated)
FileSize: 62.9 MB (estimated)
WARNING: Snappy compression is not enabled
```

```
-----
fillseq      :      1.197 micros/op;   92.4 MB/s
```

	Level	Files	Size(MB)	Time(sec)	Read(MB)	Write(MB)
--	-------	-------	----------	-----------	----------	-----------

	2	33	102	0	0	108
	3	2	6	0	0	0

# Appendix

## - Batch processing

### Advantage

- Accuracy (No human errors)
- Simplicity (No special systems)
- Efficiency (Offline feature)
- Cost savings (Automation)

### Disadvantage

- Slow (relatively Stream process)




그림 출처 : <https://memgraph.com/blog/batch-processing-vs-stream-processing>

# Question

---

