

HOMEWORK

4. Practice 2

쥬용지에

arashio1111@gmail.com

Contents:

1. 과제 설명

2. 기능 소개

3. 결과 분석

4. Practice 2

```
[Load] $ ./db_bench --benchmarks="fillrandom" --use_existing_db=0
```

```
[A] $ ./db_bench --benchmarks="readseq" --use_existing_db=1
```

```
[B] $ ./db_bench --benchmarks="readrandom" --use_existing_db=1
```

```
[C] $ ./db_bench --benchmarks="seekrandom" --use_existing_db=1
```

Note - Before running A, B, and C, run db_load benchmark.

Q1. Which user key-value interface does each benchmark use? (Put, Get, Iterator, ...)

Q2. Compare throughput and latency of each benchmark and explain why.

Q1. Which user key-value interface does each benchmark use?

\$./db_bench --benchmarks="fillrandom" --use_existing_db=0

```
arashio@arashio:~/leveldb_release/build$ ./db_bench --benchmarks="fillrandom" --use_existing_db=0
LevelDB:    version 1.23
Date:       Mon Jul 18 04:03:52 2022
CPU:        4 * Intel(R) Core(TM) i5-4690 CPU @ 3.50GHz
CPUCache:   6144 KB
Keys:       16 bytes each
Values:     100 bytes each (50 bytes after compression)
Entries:    1000000
RawSize:    110.6 MB (estimated)
FileSize:   62.9 MB (estimated)
WARNING: Snappy compression is not enabled
-----
fillrandom  :      2.683 micros/op;   41.2 MB/s
```

Q1. Which user key-value interface does each benchmark use?


```
$ ./db_bench --benchmarks="readrandom" --use_existing_db=1
```

```
void ReadRandom(ThreadState* thread) {
    ReadOptions options;
    std::string value;
    int found = 0;
    KeyBuffer key;
    for (int i = 0; i < reads_; i++) {
        const int k = thread->rand.Uniform(FLAGS_num);
        key.Set(k);
        if (db_>Get(options, key.slice(), &value).ok()) {
            found++;
        }
        thread->stats.FinishedSingleOp();
    }
    char msg[100];
    std::snprintf(msg, sizeof(msg), "(%d of %d found)", found, num_);
    thread->stats.AddMessage(msg);
}
```

Q1. Which user key-value interface does each benchmark use?

Readrandom
= **get** (random Key) * N

Random read




Q1. Which user key-value interface does each benchmark use?

\$./db_bench --benchmarks="seekrandom" --use_existing_db=1


```
void SeekRandom(ThreadState* thread) {
    ReadOptions options;
    int found = 0;
    KeyBuffer key;
    for (int i = 0; i < reads_; i++) {
        Iterator* iter = db_>NewIterator(options);
        const int k = thread->rand.Uniform(FLAGS_num);
        key.Set(k);
        iter->Seek(key.slice());
        if (iter->Valid() && iter->key() == key.slice()) found++;
        delete iter;
        thread->stats.FinishedSingleOp();
    }
    char msg[100];
    snprintf(msg, sizeof(msg), "(%d of %d found)", found, num_);
    thread->stats.AddMessage(msg);
}
```

Q1. Which user key-value interface does each benchmark use?

Seekrandom
= **seek** (random Key) * N



Random read



Q1. Which user key-value interface does each benchmark use?

\$./db_bench --benchmarks="readseq" --use_existing_db=1


```
void ReadSequential(ThreadState* thread) {  
    Iterator* iter = db_>NewIterator(ReadOptions());  
    int i = 0;  
    int64_t bytes = 0;  
    for (iter->SeekToFirst(); i < reads_ && iter->Valid(); iter->Next()) {  
        bytes += iter->key().size() + iter->value().size();  
        thread->stats.FinishedSingleOp();  
        ++i;  
    }  
    delete iter;  
    thread->stats.AddBytes(bytes);  
}
```


Q1. Which user key-value interface does each benchmark use?


Readseq

= **Iterator**->SeekToFirst()
Iterator->next() *N

Sequential read



Wenshao Zhong, REMIX: Efficient Range Query for LSM-trees, FAST '21



Q2. Compare latency of each benchmark and explain why.

	read seq	read random	seek random
User interface	Iterator next()	Get()	Iterator seek()
I/O	sequential read	random read	random read
Latency	0.102 micros/op	1.482 micros/op	1.742 micros/op

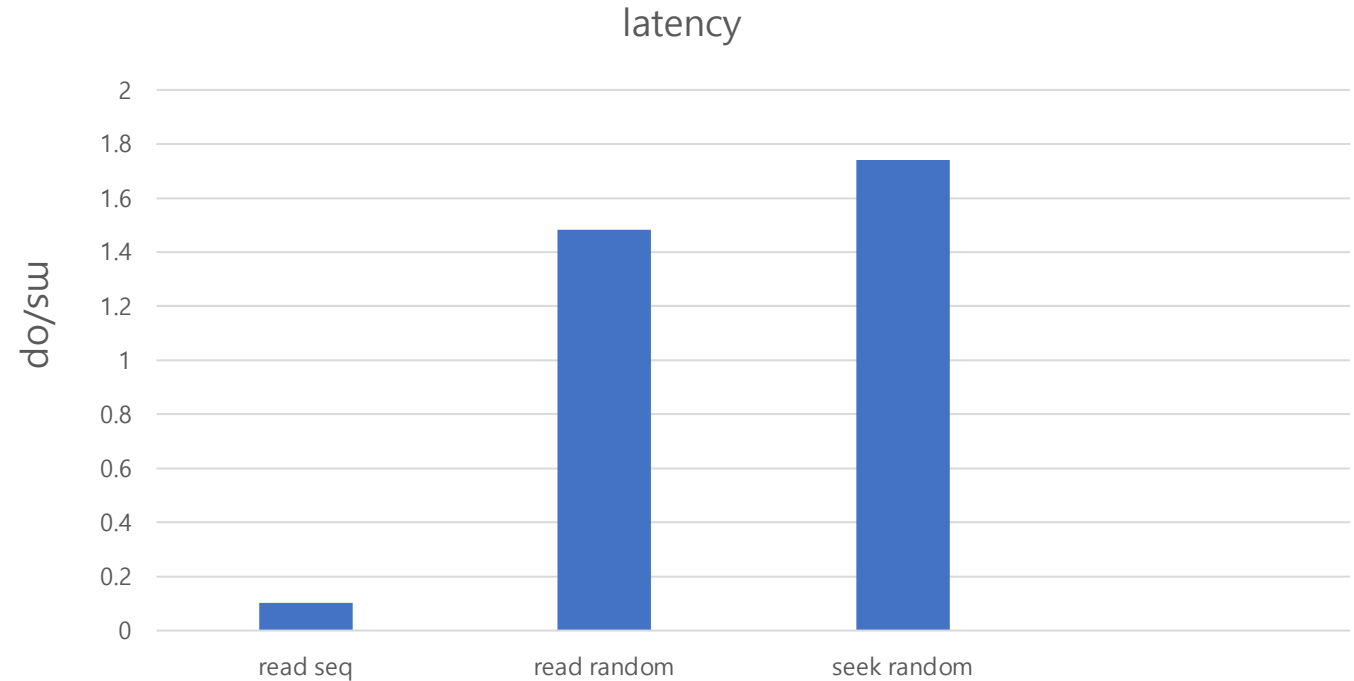
Q2. Compare latency of each benchmark and explain why.

Readrandom and **seekrandom** are random reads.

- ✓ seekrandom needs to query the highest level every time.
- ✓ seekrandom will be **slower** than readrandom.

Readseq is read sequentially

AS SSD Benchmark 1.9.5986.35387	
File Edit View Tools Language Help	
E: Samsung SSD 860 PRO 512GB 1 GB	
Samsung SSD 860 RVM01B6Q storahci - OK 1024 K - OK 476.94 GB	Read:
<input checked="" type="checkbox"/> Seq	528.27 MB/s
<input checked="" type="checkbox"/> 4K	45.47 MB/s



Q&A

감사합니다
Thank you~!

