

LevelDB Study

Introduction 3

2022. 07. 19

Presented by Min-guk Choi

koreachoi96@gmail.com

Previous Homework



leveldb 멀티 레벨 구조



why leveldb multi level

The Log-Structured Merge-Tree (LSM-Tree)

Patrick O'Neil¹, Edward Cheng²
Dieter Gawlick³, Elizabeth O'Neil¹
To be published: Acta Informatica

ABSTRACT. High-performance transaction system applications typically insert rows in a History table to provide an activity trace; at the same time the transaction system generates log records for purposes of system recovery. Both types of generated information can benefit from efficient indexing. An example in a well-known setting is the TPC-A benchmark application, modified to support efficient queries on the History for account activity for specific accounts. This requires an index by account-id on the fast-growing History table. Unfortunately, standard disk-based index structures such as the B-tree will effectively double the I/O cost of the transaction to maintain an index such as this in real time, increasing the total system cost up to fifty percent. Clearly a method for maintaining a real-time index at low cost is desirable. The Log-Structured Merge-tree (LSM-tree) is a disk-based data structure designed to provide low-cost indexing for a file experiencing a high rate of record inserts (and deletes) over an extended period. The LSM-tree uses an algorithm that defers and batches index changes, cas-

```
void InternalKeyComparator::FindShortSuccessor(std::string* key) const {
    // ...
    user_comparator->Compare(user_start, tmp) < 0) {
        // User key has become shorter physically, but larger logically.
        // Tack on the earliest possible number to the shortened user key.
        PutFixed64(&tmp,
            PackSequenceAndType(kMaxSequenceNumber, kValueTypeForSeek));
        assert(this->Compare(*start, tmp) < 0);
        assert(this->Compare(tmp, limit) < 0);
        start->swap(tmp);
    }
}

void InternalKeyComparator::FindShortSuccessor(std::string* key) const {
    Slice user_key = ExtractUserKey(*key);
    std::string tmp(user_key.data(), user_key.size());
    user_comparator->FindShortSuccessor(&tmp);
    if (tmp.size() < user_key.size() &&
        user_comparator->Compare(user_key, tmp) < 0) {
        // User key has become shorter physically, but larger logically.
        // Tack on the earliest possible number to the shortened user key.
    }
}
```

```
mingu@sever:~/leveldb_release/build$ ./db_bench
LevelDB:  version 1.23
Date:      Fri Jul  8 21:38:00 2022
CPU:       16 * Intel(R) Core(TM) i7-10700K CPU @ 3.80GHz
CPUCache:  16384 KB
Keys:      16 bytes each
Values:    100 bytes each (50 bytes after compression)
Entries:   1000000
RawSize:   110.6 MB (estimated)
FileSize:  62.9 MB (estimated)
```

1. *How to analyze*

LevelDB

2. VS code

3. Understand

4. GDB

5. Uftrace

■ How to analyze LevelDB

✓ How to analyze LevelDB

✓ How to use

✓ What/When to use?

✓ Where to start?

✓ Source Code

How to analyze LevelDB

1. Do a research
 - Documents, lectures.
2. Remarks, Code
 - VS code, Understand
3. Code Tracing
 - GDB, Uftrace
4. Draw figures
 - Structure, class, flow
 - Draw.io, PPT
5. Write a markdown document
6. Prepare a 15-minute presentation

How I use

File Browsing

Searching

VS code

- Go to Definition

- Go to References

uftrace tui

Gdb

- Go to code

db_bench

Understand

- Butterfly
- Object references
- Class Diagram

Draw.io




Paper

What/When to use?

	1	2
Static Analysis Tool - Remarks & Code	VS code - Go to Definition - Go to References - Search - File Explorer	Understand - Class Diagram - Object References - Butterfly
Dynamic Analysis Tool - Code Flow/Tracing	Uftrace - Record - Replay - Tracing	GDB - Break point - Step into / Next - Print / Display - Line by Line - Argument, Variable

Where to start?

- Start with basic operations
 - Put(Write) operations
 - db_bench --benchmarks="fillseq, fillrandom"
 - db_bench.cc:DoWrite
 - Get(Read) operations
 - db_bench --benchmarks="readrandom"
 - db_bench.cc:ReadRandom
 - Seek(Scan) operations
 - db_bench --benchmarks="readseq"
 - db_bench.cc:ReadSequential



Source Code

Subject	files
WAL/Manifest	version_set.h version_edit.h write_batch.cc db_impl.h Repair.cc
MemTable	Skiplist.h memtable.h db_impl.h Arena.h
Compaction	db_impl.h

Subject	files
SSTable	table/
Bloom Filter	c.cc Dbformat.cc filter_block.cc filter_block_test.cc filter_policy.h
Cache	cache.h table.cc table_cache.h hash.h db_impl.h

1. How to analyze

LevelDB

2. *VS code*

3. Understand

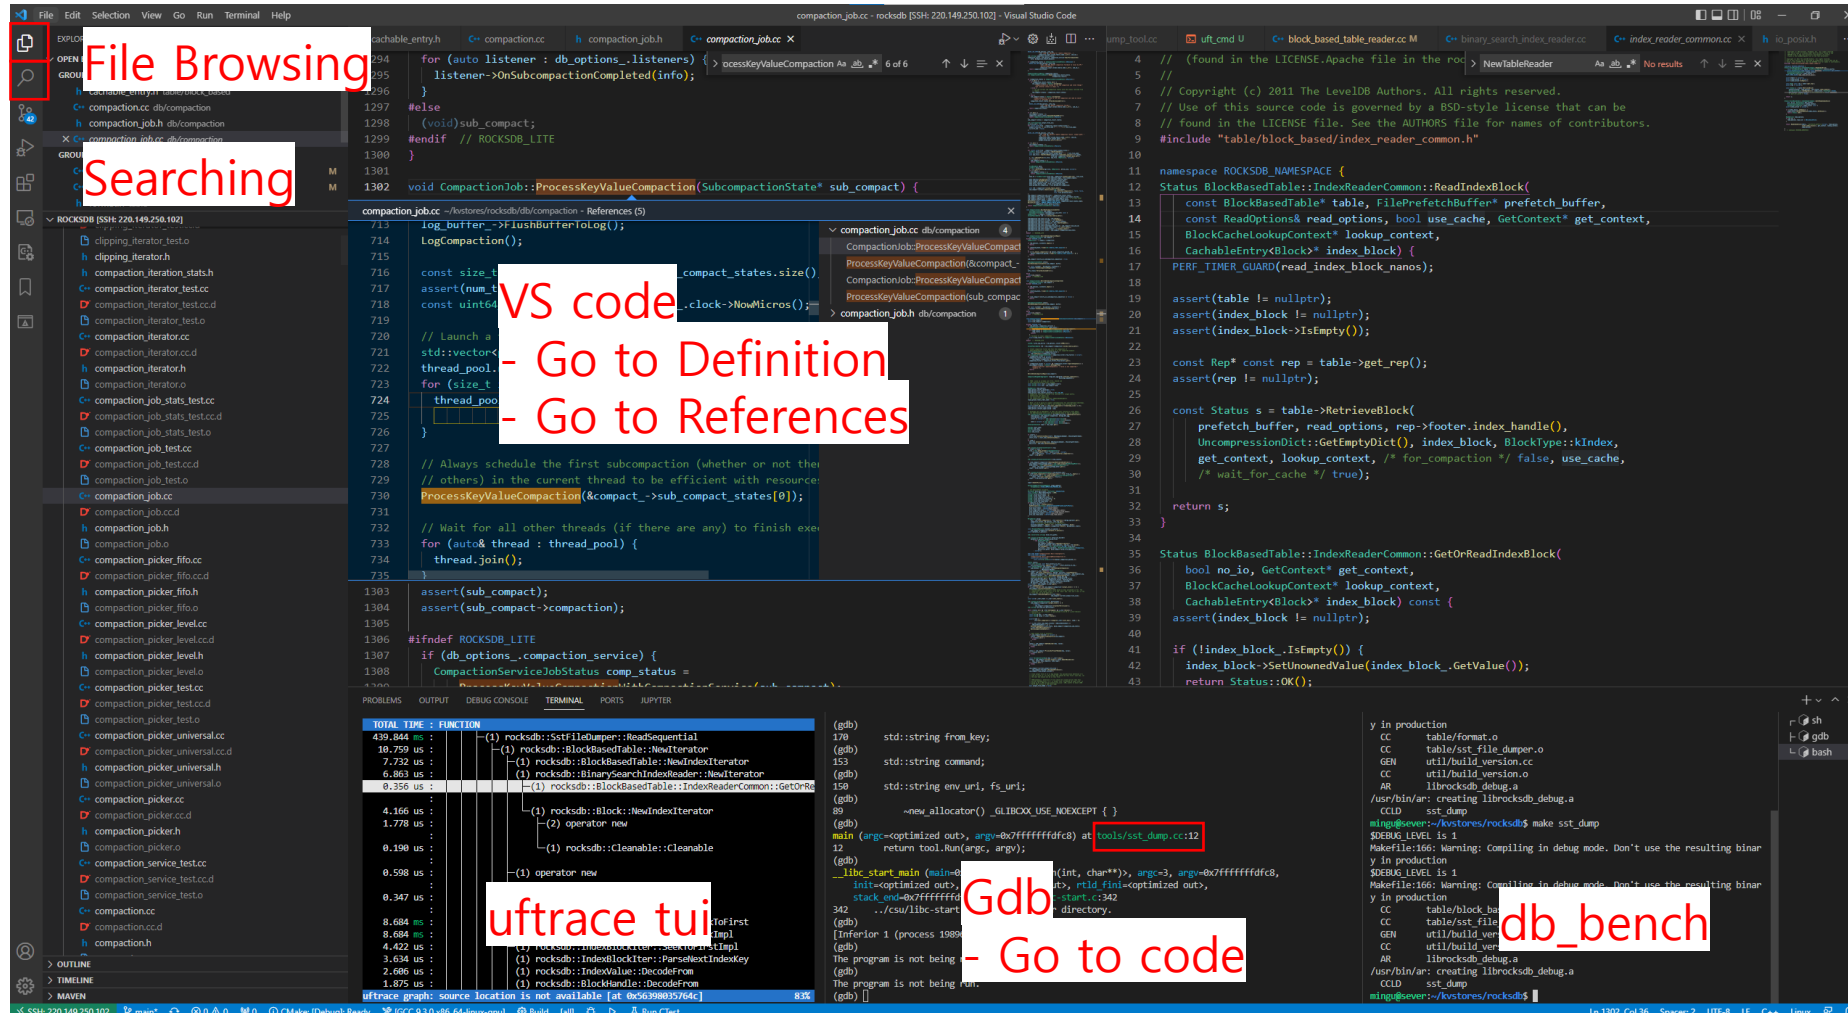
4. GDB

5. Uftrace

- VS Code

- ✓ References

VS Code



VS Code

- References

- 25 VS Code Productivity Tips and Speed Hacks
 - <https://youtu.be/ifTF3ags0XI>
- 코딩시간을 절반으로 줄여주는 VSCode 9개 기능
 - <https://youtu.be/mh-0twurNRE>
- 비주얼 스튜디오 코드 필수 단축키 정리 (Visual Studio Code 꿀템 🍯 🐼)
 - <https://youtu.be/EVxCdenPbFs>
- Visual Studio Code 기본 사용법
 - <https://youtu.be/K8qVH8V0VvY>

1. How to analyze

LevelDB

2. VS code

3. *Understand*

4. GDB

5. Uftrace

- Understand

- ✓ Introduction

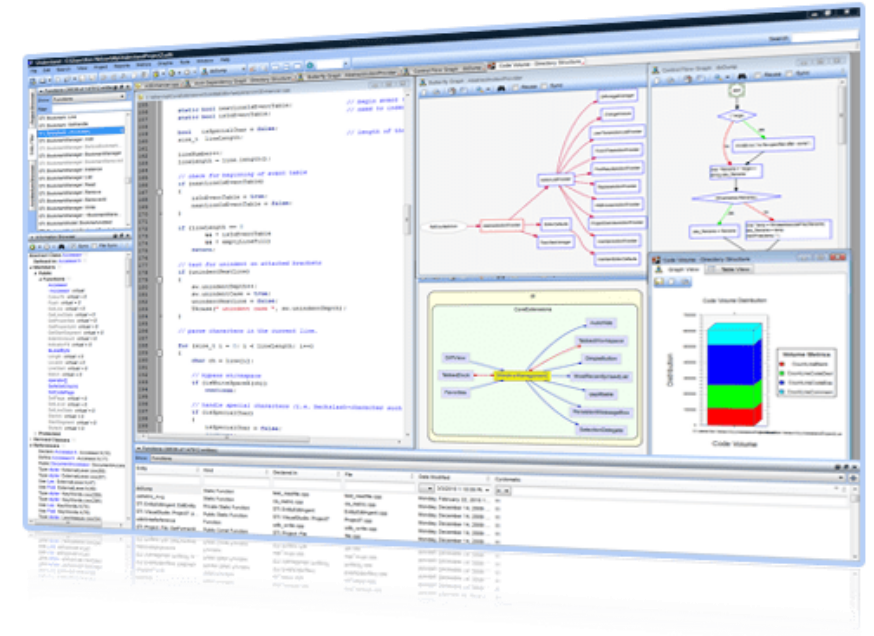
- ✓ Features

- Class diagram
 - Object references
 - Control flow
 - Butterfly

- ✓ Installation


Understand

- Understand by SciTools
 - <https://www.scitools.com/> (ENG)
 - <https://www.slexn.com/understand/> (KOR)
- Static analysis tool
 - understanding of complex open-source code
 - Graphing
 - code dependencies, code flows, function calls, and more

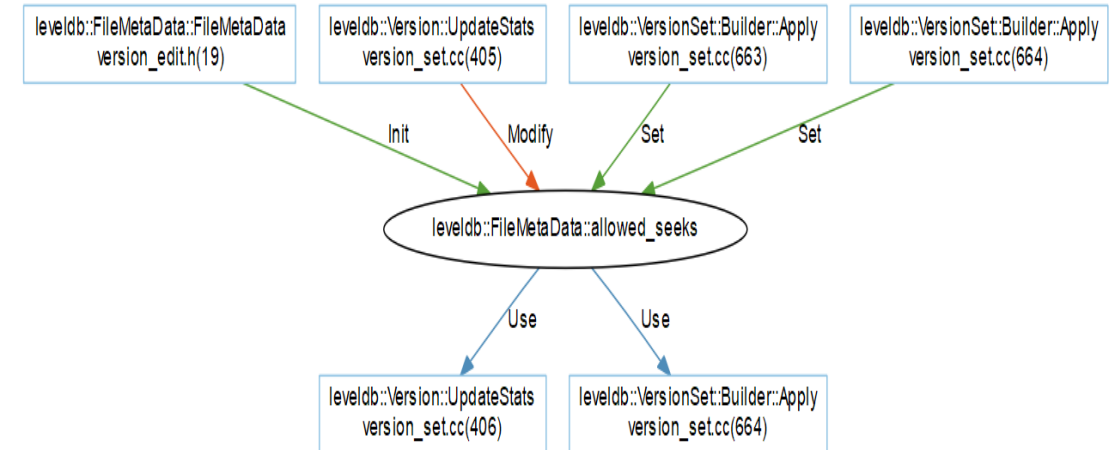


Understand

- LevelDB example




UML Class Diagram: Iterator



Object References

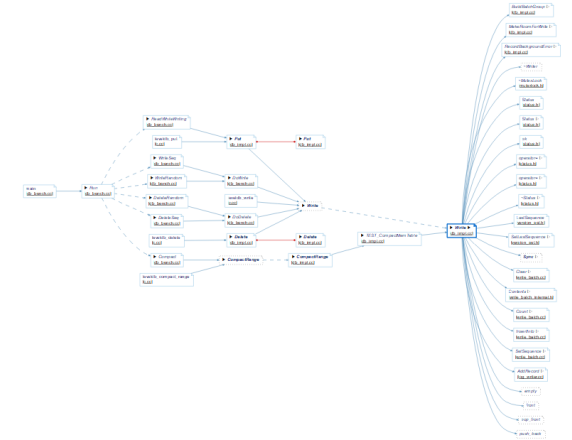


Understand

- LevelDB example
 - Control flow: db_impl.cc:Write
 - suggest to read code using VS code
 - Use “Go to Definition” (F12)




Understand

- LevelDB example
 - Butterfly: db_impl.cc:Get
 - Calls, called by



Understand

- How to install
 - Register with university email
 - <https://licensing.scitools.com/register>
 - Free trial link
 - <https://www.scitools.com/pricing>
 - Download link
 - <https://licensing.scitools.com/download>



Free for Educational Use

Understand is free for students and teachers to use for educational purposes. Want to teach a class on code maintenance or need some metrics for your thesis? No problem, we got you covered. [Learn More](#)

Understand

- How to use Understand?
 - <https://support.scitools.com/support/solutions>

General (12)

- 📖 Does Understand run on Windows 11?
- 📖 Maximizing Performance on a Large Code Base
- 📖 Using Understand from the Command Line ...
- 📖 Information Browser
- 📖 Git Integration

[View all 12](#)

Videos (19)

- 📖 Context Menu Video
- 📖 Graph Overview Video
- 📖 Architecture Designer Video
- 📖 Introduction to Understand
- 📖 Creating a New Project From Source

[View all 19](#)

Graphs (8)

- 📖 Graph for Assignments and Assigned By
- 📖 What's Changed with Graphs in 6.1?
- 📖 UML Class Diagram
- 📖 UML Sequence Diagram
- 📖 Cluster Graph Styles

[View all 8](#)

1. How to analyze

LevelDB

2. VS code

3. Understand

4. ***GDB***

5. Uftrace


- GDB

- ✓ Introduction


- ✓ Command

- ✓ Example

- Compaction Merge Iterator



- What is GDB?
 - GNU Project debugger
 - see what is going on 'inside' another program while it executes
 - see what another program was doing the moment it crashed.
- Key features
 - Start your program, specifying anything that might affect its behavior.
 - Make your program stop on specified conditions.
 - Examine what has happened when your program has stopped.
 - Change things in your program so you can experiment with correcting the effects of one bug and go on to learn about another.



GNU

gdb

- Run

\$ gdb <program>

\$ gdb --args <program> <arg1> <arg2> ...

- Break Point

>b <file_name>:<line>

>i b

>d <breakpoint #>

>en <breakpoint #>

>dis <breakpoint #>

> b <file_name>:<function>

>info break

>delete <breakpoint #>

>enable <breakpoint #>

>disable <breakpoint #>

gdb

- **Process**

<i>>r</i>	<i>>run</i>
<i>>c</i>	<i>>continue</i>
<i>>n</i>	<i>>next</i>
<i>>s</i>	<i>>step</i>
<i>>fin</i>	<i>>finish</i>

- **Print**

<i>>p <val></i>	<i>>print <val></i>
<i>>display <val></i>	
<i>>info locals</i>	
<i>>info variables</i>	
<i>>info f</i>	

LevelDB Example: Merge Iterator

```
mingu@server:~/leveldb_release/build$ gdb --args ./db_bench --benchmarks="fillrandom"
```

```
GNU gdb (Ubuntu 9.2-0ubuntu1~20.04.1) 9.2
```

```
Copyright (C) 2020 Free Software Foundation, Inc.
```

```
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
```

```
This is free software: you are free to change and redistribute it.
```

```
There is NO WARRANTY, to the extent permitted by law.
```

```
Type "show copying" and "show warranty" for details.
```

```
This GDB was configured as "x86_64-linux-gnu".
```

```
Type "show configuration" for configuration details.
```

```
For bug reporting instructions, please see:
```

```
<http://www.gnu.org/software/gdb/bugs/>.
```

```
Find the GDB manual and other documentation resources online at:
```

```
<http://www.gnu.org/software/gdb/documentation/>.
```

```
For help, type "help".
```

```
Type "apropos word" to search for commands related to "word"...
```

```
Reading symbols from ./db_bench...
```

Open file in Editor (Ctrl + Click)

```
(gdb) b db_impl.cc:894
```

```
Breakpoint 1 at 0x12213: file /home/mingu/leveldb_release/db/db_impl.cc, line 894.
```

```
(gdb) i b
```

Num	Type	Disp	Enb	Address	What
1	breakpoint	keep y		0x00000000000012213	in leveldb::DBImpl::DoCompactionWork(leveldb::DBImpl::CompactionState*)

```
at /home/mingu/leveldb_release/db/db_impl.cc:894
```

```
(gdb) █
```

LevelDB Example: Merge Iterator

```
(gdb) r
Starting program: /home/mingu/leveldb_release/build/db_bench --benchmarks=fillrandom
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
LevelDB:    version 1.23
Date:       Fri Jul 15 15:57:24 2022
CPU:        16 * Intel(R) Core(TM) i7-10700K CPU @ 3.80GHz
CPUCache:   16384 KB
Keys:       16 bytes each
Values:     100 bytes each (50 bytes after compression)
Entries:    1000000
RawSize:    110.6 MB (estimated)
FileSize:   62.9 MB (estimated)
WARNING: Optimization is disabled: benchmarks unnecessarily slow
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
[New Thread 0x7ffff79bd700 (LWP 7544)]
[New Thread 0x7ffff6eca700 (LWP 7545)]
[Switching to Thread 0x7ffff6eca700 (LWP 7545)]

Thread 3 "db_bench" hit Breakpoint 1, leveldb::DBImpl::DoCompactionWork (this=0x5555555dbc70, compact=0x7ffffe8000c20)
    at /home/mingu/leveldb_release/db/db_impl.cc:894
894      int64_t imm_micros = 0; // Micros spent doing imm_ compactions
(gdb) |
```


LevelDB Example: Merge Iterator

```
1013     input->Next();
(gdb) s
leveldb::(anonymous namespace)::MergingIterator::Next (this=
    0x555555590fee <leveldb::(anonymous namespace)::TwoLevelIterator::value() const+96>)
    at /home/mingu/leveldb_release/table/merger.cc:55
55     void Next() override {
(gdb) n
56     assert(Valid());
(gdb)
63     if (direction_ != kForward) {
(gdb)
77     current_->Next();
(gdb)
78     FindSmallest();
(gdb) s
leveldb::(anonymous namespace)::MergingIterator::FindSmallest (
    this=0x555555590f8c <leveldb::(anonymous namespace)::TwoLevelIterator::key() const+96>)
    at /home/mingu/leveldb_release/table/merger.cc:148
148     void MergingIterator::FindSmallest() {
(gdb)
```

LevelDB Example: Merge Iterator

```
(gdb) display smallest
1: smallest = (leveldb::IteratorWrapper *) 0x0
(gdb) display child
2: child = (leveldb::IteratorWrapper *) 0x7fffe8010438
(gdb) n
150     for (int i = 0; i < n_; i++) {
1: smallest = (leveldb::IteratorWrapper *) 0x7fffe8010438
(gdb)
151     IteratorWrapper* child = &children_[i];
1: smallest = (leveldb::IteratorWrapper *) 0x7fffe8010438
2: child = (leveldb::IteratorWrapper *) 0x7fffe8010438
(gdb) display *smallest
3: *smallest = {iter_ = 0x7fffe8012110, valid_ = true, key_ = {
    data_ = 0x7fffe8012b70 '0' <repeats 13 times>, "198\001\225X\002", size_ = 24}}
(gdb) display *child
4: *child = {iter_ = 0x7fffe8012110, valid_ = true, key_ = {
    data_ = 0x7fffe8012b70 '0' <repeats 13 times>, "198\001\225X\002", size_ = 24}}
(gdb) n
152     if (child->Valid()) {
1: smallest = (leveldb::IteratorWrapper *) 0x7fffe8010438
2: child = (leveldb::IteratorWrapper *) 0x7fffe8010458
3: *smallest = {iter_ = 0x7fffe8012110, valid_ = true, key_ = {
    data_ = 0x7fffe8012b70 '0' <repeats 13 times>, "198\001\225X\002", size_ = 24}}
4: *child = {iter_ = 0x7fffe800ff50, valid_ = true, key_ = {
    data_ = 0x7fffe8012e30 '0' <repeats 14 times>, "40\001\213\206\001", size_ = 24}}
(gdb)
```

```
void MergingIterator::FindSmallest() {
    IteratorWrapper* smallest = nullptr;
    for (int i = 0; i < n_; i++) {
        IteratorWrapper* child = &children_[i];
        if (child->Valid()) {
            if (smallest == nullptr) {
                smallest = child;
            } else if (comparator_->Compare(child->key(),
                smallest->key()) < 0) {
                smallest = child;
            }
        }
    }
    current_ = smallest;
}
```

LevelDB Example: Merge Iterator

```
void MergingIterator::FindSmallest() {
    IteratorWrapper* smallest = nullptr;
    for (int i = 0; i < n_; i++) {
        IteratorWrapper* child = &children_[i];
        if (child->Valid()) {
            if (smallest == nullptr) {
                smallest = child;
            } else if (comparator_>Compare(child->key(),
                smallest->key()) < 0) {
                smallest = child;
            }
        }
    }
    current_ = smallest;
}
```

```
156     smallest = child;
3: *smallest = {iter_ = 0x7fffe800ff50, valid_ = true, key_ = {
    data_ = 0x7fffe8012e30 '0' <repeats 14 times>, "40\001\213\206\001", size_ = 24}}
4: *child = {iter_ = 0x7fffe80100a0, valid_ = true, key_ = {
    data_ = 0x7fffe800fa30 '0' <repeats 14 times>, "30\001\211\024\002", size_ = 24}}
(gdb)
150     for (int i = 0; i < n_; i++) {
3: *smallest = {iter_ = 0x7fffe80100a0, valid_ = true, key_ = {
    data_ = 0x7fffe800fa30 '0' <repeats 14 times>, "30\001\211\024\002", size_ = 24}}
```

LevelDB Example: Merge Iterator

```
void MergingIterator::FindSmallest() {
    IteratorWrapper* smallest = nullptr;
    for (int i = 0; i < n_; i++) {
        IteratorWrapper* child = &children_[i];
        if (child->Valid()) {
            if (smallest == nullptr) {
                smallest = child;
            } else if (comparator_->Compare(child->key(),
                smallest->key()) < 0) {
                smallest = child;
            }
        }
    }
    current_ = smallest;
}
```

```
160     current_ = smallest;
3: *smallest = {iter_ = 0x7fffe80100a0, valid_ = true, key_ = {
    data = 0x7fffe800fa30 '0' <repeats 14 times>, "30\001\211\024\002", size_ = 24}}
7: *current_ = {iter_ = 0x7fffe800ff50, valid_ = true, key_ = {
    data_ = 0x7fffe8012e30 '0' <repeats 14 times>, "40\001\213\206\001", size_ = 24}}
(gdb)
161     }
3: *smallest = {iter_ = 0x7fffe80100a0, valid_ = true, key_ = {
    data = 0x7fffe800fa30 '0' <repeats 14 times>, "30\001\211\024\002", size_ = 24}}
7: *current_ = {iter_ = 0x7fffe80100a0, valid_ = true, key_ = {
    data_ = 0x7fffe800fa30 '0' <repeats 14 times>, "30\001\211\024\002", size_ = 24}}
```

References

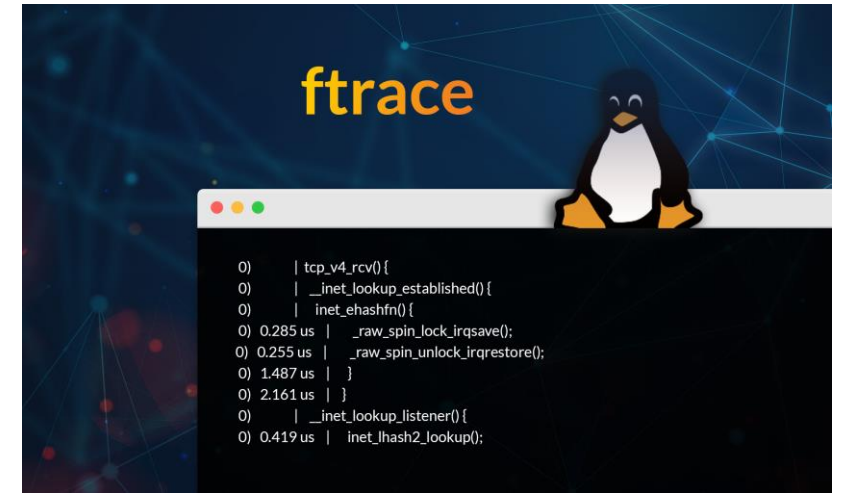
- KOR
 - gdb 디버거 사용법 및 다양한 기능 설명
 - <https://edward0im.github.io/technology/2020/09/29/gdb/>
 - gdb - 간단한 명령어/사용법/단축어 정리(cheat sheet)
 - <https://dining-developer.tistory.com/13>

1. How to analyze
LevelDB
2. VS code
3. Understand
4. GDB
5. *Uftrace*

- Uftrace
 - ✓ Features
 - Record
 - Replay
 - Filter
 - TUI
 - ✓ Shell Script
 - ✓ Installation

uftrace

- Trace and analyze execution of a program written in C/C++
- Heavily inspired by the ftrace framework of the Linux kernel
- Supports user-space programs and kernel
- Various kind of commands and filters



Features

- Record
 - Runs a program and saves the trace data
- Replay
 - Show program execution in the trace data
- Graph/Tui
 - Show function call graph in the trace data
- Filter

Replay: Write()

#	DURATION	TID	FUNCTION
		[5471]	leveldb::DBImpl::Write() {
		[5471]	leveldb::DBImpl::Writer::Writer()
0.030	us	[5471]	leveldb::Status::Status();
		[5471]	leveldb::port::CondVar::CondVar()
0.164	us	[5471]	std::condition_variable::condition_variable()
0.296	us	[5471]	} /* leveldb::port::CondVar::CondVar()
0.508	us	[5471]	} /* leveldb::DBImpl::Writer::Writer()
		[5471]	leveldb::MutexLock::MutexLock() {
		[5471]	leveldb::port::Mutex::Lock() {
		[5471]	std::mutex::lock() {
		[5471]	__gthread_mutex_lock() {
0.040	us	[5471]	__gthread_active_p();
0.313	us	[5471]	pthread_mutex_lock();
1.758	us	[5471]	} /* __gthread_mutex_lock *
2.058	us	[5471]	} /* std::mutex::lock */
2.133	us	[5471]	} /* leveldb::port::Mutex::Lock
2.217	us	[5471]	} /* leveldb::MutexLock::MutexLock
		[5471]	std::deque::push_back() {
0.029	us	[5471]	std::move();
		[5471]	std::deque::emplace_back() {
0.031	us	[5471]	std::forward();
		[5471]	std::allocator_traits::construct
0.029	us	[5471]	std::forward();

- Record

\$ ufttrace record ./db_bench benchmarks="fillrandom" \
--num=100000

- Replay

\$ ufttrace replay

> /leveldb::DBImpl::Write

\$ ufttrace replay -F leveldb::DBImpl::Write

Filter: Write()

#	DURATION	TID	FUNCTION
		[5471]	leveldb::DBImpl::Write() {
		[5471]	leveldb::DBImpl::Writer::Writer() {
0.030 us		[5471]	leveldb::Status::Status();
		[5471]	leveldb::port::CondVar::CondVar() {
0.164 us		[5471]	std::condition_variable::condition_variable();
0.296 us		[5471]	} /* leveldb::port::CondVar::CondVar */
0.508 us		[5471]	} /* leveldb::DBImpl::Writer::Writer */
		[5471]	leveldb::MutexLock::MutexLock() {
		[5471]	leveldb::port::Mutex::Lock() {
		[5471]	std::mutex::lock() {
		[5471]	__gthread_mutex_lock() {
0.040 us		[5471]	__gthread_active_p();
0.313 us		[5471]	pthread_mutex_lock();
1.758 us		[5471]	} /* __gthread_mutex_lock */
2.058 us		[5471]	} /* std::mutex::lock */
2.133 us		[5471]	} /* leveldb::port::Mutex::Lock */
2.217 us		[5471]	} /* leveldb::MutexLock::MutexLock */
		[5471]	std::deque::push_back() {
0.029 us		[5471]	std::move();
		[5471]	std::deque::emplace_back() {
0.031 us		[5471]	std::forward();
		[5471]	std::allocator_traits::construct() {
0.029 us		[5471]	std::forward();

No Filter

#	DURATION	TID	FUNCTION
		[14234]	leveldb::DBImpl::Write() {
		[14234]	leveldb::DBImpl::Writer::Writer() {
0.311 us		[14234]	leveldb::port::CondVar::CondVar();
0.493 us		[14234]	} /* leveldb::DBImpl::Writer::Writer */
1.963 us		[14234]	leveldb::MutexLock::MutexLock();
		[14234]	leveldb::DBImpl::MakeRoomForWrite() {
0.262 us		[14234]	leveldb::VersionSet::NumLevelFiles();
0.651 us		[14234]	leveldb::MemTable::ApproximateMemoryUsage();
1.637 us		[14234]	} /* leveldb::DBImpl::MakeRoomForWrite */
0.079 us		[14234]	leveldb::VersionSet::LastSequence();
		[14234]	leveldb::DBImpl::BuildBatchGroup() {
0.161 us		[14234]	leveldb::WriteBatchInternal::ByteSize();
1.423 us		[14234]	} /* leveldb::DBImpl::BuildBatchGroup */
0.315 us		[14234]	leveldb::WriteBatchInternal::SetSequence();
0.221 us		[14234]	leveldb::WriteBatchInternal::Count();
0.328 us		[14234]	leveldb::WriteBatchInternal::Contents();
		[14234]	leveldb::log::Writer::AddRecord() {
		[14234]	leveldb::log::Writer::EmitPhysicalRecord() {
0.757 us		[14234]	leveldb::_GLOBAL__N_1::PosixWritableFile::Append();
0.578 us		[14234]	leveldb::_GLOBAL__N_1::PosixWritableFile::Append();

Filter

Filter

COMMON OPTIONS

-F *FUNC*, --filter=*FUNC* : Set filter to trace selected functions and their children functions. This option can be used more than once. See *FILTERS*.

-N *FUNC*, --notrace=*FUNC* : Set filter not to trace selected functions and their children functions. This option can be used more than once. See *FILTERS*.

-C *FUNC*, --caller-filter=*FUNC* : Set filter to trace callers of selected functions only. This option can be used more than once. See *FILTERS*.

-T *TRG*, --trigger=*TRG* : Set trigger on selected functions. This option can be used more than once. See *TRIGGERS*.

-D *DEPTH*, --depth *DEPTH* : Set trace limit in nesting level. See *FILTERS*.

-t *TIME*, --time-filter=*TIME* : Do not show functions which run under the time threshold. If some functions explicitly have the 'trace' trigger applied, those are always traced regardless of execution time. See *FILTERS*.

--no-libcall : Do not show library calls.

COMMON ANALYSIS OPTIONS

-H *FUNC*, --hide=*FUNC* : Set filter not to trace selected functions. It doesn't affect their subtrees, but hides only the given functions. This option can be used more than once. See *FILTERS*.

--kernel-full : Show all kernel functions and events occurred outside of user functions.

--kernel-only : Show kernel functions only without user functions.

--event-full : Show all (user) events outside of user functions.

--tid=*TID*[,*TID*,...] : Only print functions called by the given tasks. To see the list of tasks in the data file, you can use `uftrace report --task` or `uftrace info`. This option can also be used more than once.

[Uftrace/doc/uftrace-replay.md](https://uftrace.org/doc/uftrace-replay.md)

Tui: Write()

TOTAL TIME : FUNCTION	
11.538 s	(1) db_bench
5.502 s	(85352) leveldb::DBImpl::Write
26.224 ms	(85352) leveldb::DBImpl::Writer::Writer
12.326 ms	(85352) leveldb::port::CondVar::CondVar
39.723 ms	(85352) leveldb::MutexLock::MutexLock
72.355 ms	(85352) leveldb::DBImpl::MakeRoomForWrite
10.459 ms	(85358) leveldb::VersionSet::NumLevelFiles
16.385 ms	(85355) leveldb::MemTable::ApproximateMemoryUsage
0.090 us	(3) leveldb::VersionSet::PrevLogNumber
0.100 us	(3) leveldb::VersionSet::NewFileNumber
6.257 us	(3) leveldb::LogFileName
5.965 us	(3) leveldb::MakeFileName
71.581 us	(3) leveldb::_GLOBAL__N_1::PosixEnv::NewWritableFile
8.899 us	(3) leveldb::_GLOBAL__N_1::PosixWritableFile::PosixWritableFile
0.103 us	(3) leveldb::WritableFile::WritableFile
4.749 us	(3) leveldb::_GLOBAL__N_1::PosixWritableFile::IsManifest
2.985 us	(3) leveldb::_GLOBAL__N_1::PosixWritableFile::Basename
2.038 us	(3) leveldb::_GLOBAL__N_1::PosixWritableFile::Dirname
0.093 us	(3) leveldb::log::Writer::~Writer
10.424 us	(3) leveldb::_GLOBAL__N_1::PosixWritableFile::~PosixWritableFile
8.303 us	(3) leveldb::_GLOBAL__N_1::PosixWritableFile::~PosixWritableFile
6.650 us	(3) leveldb::_GLOBAL__N_1::PosixWritableFile::Close
uftrace graph: session 241c4796d8967ad9 (/home/mingu/leveldb_release/build/db_bench)	

■ Tui

- Back-trace
- Call Graph

■ Run uftrace

```
$ uftrace record ./db_bench --benchmarks="fillrandom" \
--num=100000
```

```
$ uftrace tui
```

```
>/ leveldb::DBImpl::Write
```

```
>g
```

Tui: SkipList::Insert

```
===== Back-trace =====
5.188 s : (100000) leveldb::SkipList::Insert
5.188 s : (100000) leveldb::MemTable::Add
5.188 s : (100000) leveldb::_GLOBAL__N_1::MemTableInserter::Put
5.188 s : (100000) leveldb::WriteBatch::Iterate
5.188 s : (100000) leveldb::WriteBatchInternal::InsertInto
5.188 s : (100000) leveldb::DBImpl::Write
5.188 s : (100000) leveldb::Benchmark::DoWrite
5.188 s : (100000) leveldb::Benchmark::WriteRandom
5.188 s : (100000) leveldb::Benchmark::ThreadBody
:
===== Call Graph =====
5.188 s : (100000) leveldb::SkipList::Insert
4.824 s : |-(100000) leveldb::SkipList::FindGreaterOrEqual
13.450 ms : |-(100000) leveldb::SkipList::GetMaxHeight
:
473.144 ms : |-(2495692) leveldb::SkipList::Node::Next
:
4.121 s : |-(2495692) leveldb::SkipList::KeyIsAfterNode
3.903 s : |-(2404364) leveldb::MemTable::KeyComparator::operator()
993.576 ms : |-(4808728) leveldb::GetLengthPrefixedSlice
:
2.544 s : |-(2404364) leveldb::InternalKeyComparator::Compare
492.141 ms : |-(2404364) leveldb::_GLOBAL__N_1::BytewiseComparatorImpl::Compare
:
25.035 ms : |-(100000) leveldb::SkipList::RandomHeight
16.364 ms : |-(133759) leveldb::Random::OneIn
4.255 ms : |-(133759) leveldb::Random::Next
:
11.891 ms : |-(100026) leveldb::SkipList::GetMaxHeight
:
30.790 ms : |-(100000) leveldb::SkipList::NewNode
5.508 ms : |-(100000) leveldb::Arena::AllocateAligned
1.455 ms : |-(890) leveldb::Arena::AllocateFallback
1.384 ms : |-(890) leveldb::Arena::AllocateNewBlock
```

Help: (press any key to exit)

ARROW	Navigation
PgUp/Dn	
Home/End	
Enter	Fold/unfold graph or Select session
G	Show (full) call graph
g	Show call graph for this function
R	Show uftrace report
r	Show uftrace report for this function
s	Sort by the next column in report
I	Show uftrace info
S	Change session
O	Open editor
c/e	Collapse/Expand direct children graph
C/E	Collapse/Expand all descendant graph
n/p	Next/Prev sibling
u	Move up to parent
l	Move to the longest executed child
j/k	Move down/up
z	Set current line to the center of screen
/	Search
</>/N/P	Search next/prev
v	Show debug message
f	Customize fields in graph or report mode
h/?	Show this help
q	Quit

Shell script

```
1  #!/bin/sh
2
3  default_func=""
4
5  if [ -z "$1" ]
6  then
7      echo '[USAGE] : sh uft_cmd [COMMAND] [TARGET FUNC]'
8      echo '[COMMAND] : replay, report, graph, tui'
9      exit
10 fi
11
12 if [ -n "$2" ]
13 then
14     func="-F ${2}"
15 else
16     if [ -z "$default_func" ]
17     then
18         func=""
19     else
20         func="-F ${default_func}"
21     fi
22 fi
23
24 ufttrace $1 \
25 $func \
26 --no-libcall \
27 -N leveldb::MutexLock \
28 -N leveldb::ExtractUserKey \
29 -N leveldb::Arena::MemoryUsage \
30 -N __gthread_mutex_unlock \
31 -N __gthread_mutex_lock \
32 -N ^leveldb::Slice:: \
33 -N ^leveldb::port::Mutex \
34 -N ^leveldb::crc32c:: \
```

■ Shell script

- Easy to use filter
- ufttrace_script.sh

■ Run Shell script

\$ ufttrace record ./db_bench benchmarks="fillrandom" \

--num=100000

\$ sh ufttrace_script.sh replay

\$ sh ufttrace_script.sh tui

Install Guide

■ Install ufttrace

- ufttrace/INSTALL.md
 - <https://github.com/namhyung/ufttrace/blob/master/INSTALL.md>

■ Bind with leveldb

- ufttrace wiki – LevelDB/RocksDB (YCSB)
 - <https://github.com/namhyung/ufttrace/wiki/ufttrace-for-LevelDB-RocksDB-with-YCSB>

■ Tutorial

- <https://ufttrace.github.io/slide/#1>

■ Wiki

- <https://github.com/namhyung/ufttrace/wiki>

🔗 QUICK GUIDE

On Linux distros, following commands will build and install ufttrace from source.

```
$ sudo misc/install-deps.sh # optional for advanced features
$ ./configure               # --prefix can be used to change install dir
$ make
$ sudo make install
```

The screenshot shows the GitHub repository for 'namhyung / ufttrace'. The main heading is 'ufttrace for LevelDB RocksDB with YCSB'. Below the heading, it says 'Min-guk Choi edited this page 19 days ago · 6 revisions'. The page content includes a list of steps: 1. Download LevelDB, 2. Add -pg option on leveldb/CMakeList.txt for ufttrace, and 3. Build. A sidebar on the right shows a navigation menu with links to Home, Tutorial, Arguments, Filters, Development, Data Format, TODO Items, Testing, Code Coverage, Korean man page, Practical Use Cases, GCC, Clang/LVM, Node.js, Chromium, MySQL/InnoDB, FFmpeg, and CPython.

Thank you